

# BLEND: A Fast, Memory-Efficient, and Accurate Mechanism to Find Fuzzy Seed Matches

Can Firtina<sup>1</sup> Jisung Park<sup>1</sup> Mohammed Alser<sup>1</sup> Jeremie S. Kim<sup>1</sup> Damla Senol Cali<sup>2</sup>  
Taha Shahroodi<sup>3</sup> Nika Mansouri-Ghiasi<sup>1</sup> Gagandeep Singh<sup>1</sup> Konstantinos Kanellopoulos<sup>1</sup>  
Can Alkan<sup>4</sup> Onur Mutlu<sup>1</sup>

<sup>1</sup>ETH Zurich

<sup>2</sup>Bionano Genomics

<sup>3</sup>TU Delft

<sup>4</sup>Bilkent University

Generating the hash values of short subsequences, called *seeds*, enables quickly identifying similarities between genomic sequences by matching seeds with a single lookup of their hash values. However, these hash values can be used only for finding exact-matching seeds as the conventional hashing methods assign distinct hash values for different seeds, including highly similar seeds. Finding only exact-matching seeds causes either 1) increasing the use of the costly sequence alignment or 2) limited sensitivity.

We introduce BLEND, the first efficient and accurate mechanism that can identify both exact-matching and highly similar seeds with a single lookup of their hash values, called fuzzy seeds matches. BLEND 1) utilizes a technique called SimHash, that can generate the same hash value for similar sets, and 2) provides the proper mechanisms for using seeds as sets with the SimHash technique to find fuzzy seed matches efficiently.

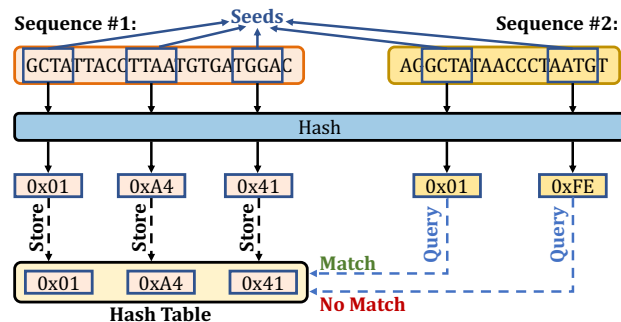
We show the benefits of BLEND when used in read overlapping and read mapping. For read overlapping, BLEND is faster by  $2.6\times$ – $63.5\times$  (on average  $19.5\times$ ), has a lower memory footprint by  $0.9\times$ – $9.7\times$  (on average  $3.6\times$ ), and finds higher quality overlaps leading to accurate de novo assemblies than the state-of-the-art tool, minimap2. For read mapping, BLEND is faster by  $0.7\times$ – $3.7\times$  (on average  $1.7\times$ ) than minimap2. Source code is available at <https://github.com/CMU-SAFARI/BLEND>.

## 1. Introduction

High-throughput sequencing (HTS) technologies have revolutionized the field of genomics due to their ability to produce millions of nucleotide sequences at a relatively low cost [1]. Although HTS technologies are key enablers of almost all genomics studies [2–7], HTS technology-provided data comes with two key shortcomings. First, HTS technologies can only provide many short (e.g., one hundred up to a million bases depending on the technology [1]) sequences (i.e., *reads*) within a full genome [8] (e.g., billions of bases for the human genome). Second, HTS technologies can misinterpret signals during sequencing and thus provide reads that contain errors (i.e., sequencing errors [8]). The average frequency of sequencing errors in a read highly varies (from 0.1% up to 15%) depending on the HTS technology [9–13]. To address the shortcomings of HTS technologies, various computational approaches must be taken to process the reads into meaningful information accurately and efficiently. These include 1) read mapping [14–18], 2) *de novo* assembly [19–21], 3) read classification in metagenomic studies [22–24], 4) correcting sequencing errors [25–27], and 5) hardware acceleration for fast sequencing data analysis [28–39].

At the core of these computational approaches, similarities between sequences must be identified to overcome the fundamental limitations of HTS technologies. However, identifying the similarities across all pairs of sequences is not practical due to the costly algorithms used to calculate the distance between two sequences, such as sequence alignment algorithms using dynamic programming (DP) approaches [40, 41]. To practically

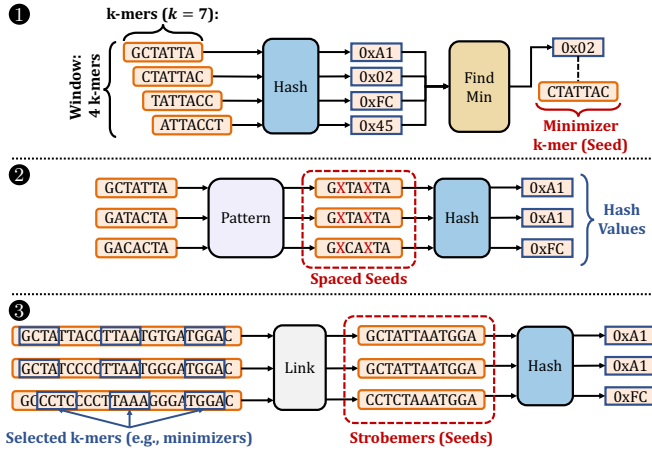
identify similarities, it is essential to avoid calculating the distance between dissimilar sequence pairs. A common heuristic is to find matching short subsequences, called *seeds*, between sequence pairs. Sequences that have no or a few seed matches are quickly filtered out from performing costly sequence alignment. There are several techniques that generate seeds from sequences, known as *seeding techniques*, for accurate and efficient filtering of dissimilar sequence pairs. To find seed matches efficiently, a common approach is to match the hash values of seeds with a single lookup using a hash table that contains the hash values of all seeds of interest. Figure 1 shows an overview of how hash tables are used to find seed matches between two sequences with a single lookup. Seeding techniques identify seeds within sequences and use hash functions to generate the hash values of these seeds. These hash values are used for filling and querying the hash table depending on the source of the seeds. Querying the hash table with hash values enables finding the positions where a seed from the second sequence appears in the first sequence, if it exists, with a single lookup. The use of seeds drastically reduces the search space from all possible sequence pairs to the sequence pairs that are likely to be similar to facilitate efficient distance calculations over a large number of sequence pairs [42–44].



**Figure 1: Finding seed matches with a single lookup of hash values.**

Figure 2 shows the three main directions that existing seeding techniques take. The first direction aims to minimize the computational overhead of using and storing seeds by selectively choosing fewer seeds from all fixed-length subsequences of reads, called *k-mers*, where the fixed-length is *k*. The existing works such as minimap2 [15], MinHash [45], Winnowmap2 [46, 47], reM<sub>u</sub>val [48], and CAS [49] use sampling techniques to choose a subset of *k*-mers from all *k*-mers of a read without significantly reducing their accuracy. For example, minimap2 uses only the *k*-mers with the minimum hash value in a window of *w* consecutive *k*-mers, known as the *minimizer* *k*-mers [43] (● in Figure 2). Such a sampling approach guarantees that one *k*-mer is sampled in each window to provide a fixed sampling ratio that can be tuned to increase the probability of matching *k*-mers between reads. Alternatively, MinHash generates many hash values from each *k*-mer of a read using many hash functions. For each hash function, only the *k*-mer with the minimum hash value is used as a seed with no win-

downing guarantees. MinHash is mainly effective for matching sequences that have similar lengths since the number of hash functions is fixed for all sequences, whereas it can generate too many seeds for shorter sequences when the sequence lengths vary greatly [14]. While these k-mer selection approaches reduce the overall number of seeds to use, all of these existing works find *only* exact-matching k-mers with a single lookup as they use hash functions with *low collision* rates to generate the hash values of these k-mers. The exact-matching requirement imposes challenges when determining the k-mer length. Longer k-mer lengths significantly decrease the probability of finding exact-matching k-mers between sequences due to genetic variations and sequencing errors. Short k-mer lengths (e.g., 8-21 bases) result in matching a large number of k-mers due to both the repetitive nature of most genomes and the high probability of finding the same short k-mer frequently in a long sequence of DNA letters [50]. Although k-mers are commonly used as seeds, a seed is a more general concept that can allow substitutions, insertions and deletions (indels) when matching short subsequences between sequence pairs.



**Figure 2: Examples of common seeding techniques.** ① Finding the minimizer k-mers. ② A spaced seeding technique. Masked characters are highlighted by X in red. ③ A simple example of the strobemers technique.

The second direction is to allow substitutions when matching k-mers by *masking* (i.e., ignoring) certain characters of k-mers and using the masked k-mers as seeds ②. Predefined *patterns* determine the fixed masking positions for all k-mers. Seeds generated from masked k-mers are known as *spaced seeds* [51]. The tools such as ZOOM! [52] and SHRiMP2 [53] use spaced seeds to improve the sensitivity when mapping short reads (i.e., Illumina paired-end reads). S-conLSH [54, 55] generates many spaced seeds from each k-mer using different masking patterns to improve the sensitivity when matching spaced seeds with locality-sensitive hashing techniques. There have been recent improvements in determining the masking patterns to improve the sensitivity of spaced seeds [56, 57]. Unfortunately, spaced seeds cannot find *any arbitrary* fuzzy matches of k-mers with a single lookup due to 1) fixed patterns that allow mismatches only at certain positions of k-mers and 2) *low collision* hashing techniques that can be used for finding *only* exact-matching spaced seeds, which are key limitations in improving the sensitivity of spaced seeds.

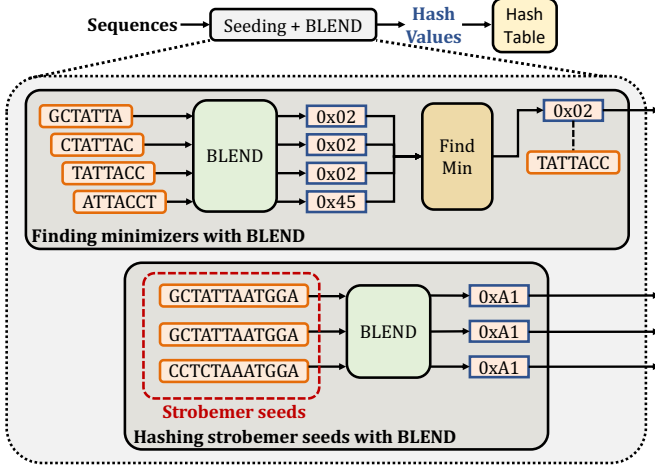
The third direction aims to allow both substitutions and indels when matching k-mers. A common approach is to link a few selected k-mers of a sequence to use these linked k-mers as seeds, such as paired-minimizers [58] and strobemers [59, 60] ③.

These approaches can ignore large gaps between the linked k-mers. For example, the strobemer technique concatenates a subset of selected k-mers of a sequence to generate a strobemer sequence, which is used as a seed. Strobemers enable masking some characters within sequences without requiring a fixed pattern, unlike spaced k-mers. This makes strobemers a more sensitive approach for detecting indels with varying lengths as well as substitutions. However, the nature of the hash function used in strobemers requires exact matches of *all* concatenated k-mers in strobemer sequences when matching seeds. Such an exact match requirement introduces challenges for further improving the sensitivity of strobemers for detecting indels and substitutions between sequences.

To our knowledge, there is no work that can *efficiently* find fuzzy matches of seeds *without* requiring 1) *exact matches* of all k-mers (i.e., any k-mer can mismatch) and 2) imposing a high performance and memory space overheads. In this work, we observe that existing works have such a limitation mainly because they employ hash functions with low collision rates when generating the hash values of seeds. Although it is important to reduce the collision rate for assigning different hash values for dissimilar seeds for accuracy and performance reasons, the choice of hash functions also makes it unlikely to assign the same hash value for similar seeds. Thus, seeds *must* exactly match to find matches between sequences with a single lookup. Mitigating such a requirement so that similar seeds can have the same hash value has the potential to improve further the performance and sensitivity of the applications that use seeds with their ability to allow substitutions and indels at any arbitrary positions when matching seeds. **Our goal** in this work is to enable finding *fuzzy* matches of seeds as well as exact-matching seeds between sequences (e.g., reads) with a single lookup of hash values of these seeds. To this end, we propose **BLEND**, a fast, memory-efficient, and accurate mechanism that can find fuzzy seed matches. The **key idea** in BLEND is to enable assigning the same hash value for highly similar seeds. To this end, BLEND 1) exploits the characteristics of the SimHash technique [61, 62] that can generate the *same* hash value for highly similar sets, and 2) provides proper mechanisms for using any seeding technique with SimHash to find fuzzy seed matches with a single lookup of their hash values. This provides us with two key benefits. First, BLEND can generate the same hash value for highly similar seeds *without* imposing exact matches of seeds, unlike existing seeding mechanisms that use hash functions with low collision rates. Second, BLEND enables finding fuzzy seed matches with a single lookup of a hash value rather than matching many hash values for calculating costly similarity scores (e.g., Jaccard similarity [63]) that the conventional locality-sensitive hashing-based methods use, such as MinHash [45] or S-conLSH [54, 55]. These two ideas ensure that BLEND can efficiently find both 1) all exact-matching seeds that a seeding technique finds using a conventional hash function with a low collision rate and 2) approximate seed matches that these conventional hashing mechanisms cannot find with a single lookup of a hash value.

Figure 3 shows two examples of how BLEND can replace the conventional hash functions that the seeding techniques use in Figure 2. The **key challenge** is to accurately and efficiently define the items of sets from seeds that the SimHash technique requires. To achieve this, BLEND provides two mechanisms for converting seeds into sets of items: 1) BLEND-I and 2) BLEND-S. To perform a sensitive detection of substitutions, BLEND-I uses all overlapping smaller k-mers of a potential seed sequence as the items of a set for generating the hash

value with SimHash. To allow mismatches between the linked k-mers that strobemers and similar seeding mechanisms use, BLEND-S uses only the linked k-mers as the set with SimHash. We envision that BLEND can be integrated with any seeding technique that uses hash values for matching seeds with a single lookup by replacing their hash function with BLEND and using the proper mechanism for converting seeds into a set of items.



**Figure 3: Replacing the hash functions in any seeding technique with BLEND.**

Using erroneous (ONT and PacBio CLR), highly accurate (PacBio HiFi), and short (Illumina) reads, we experimentally show the benefits of BLEND on two important applications in genomics: 1) read overlapping and 2) read mapping. First, read overlapping aims to find overlaps between all pairs of reads based on seed matches. These overlapping reads are mainly useful for generating an assembly of the sequenced genome [14, 64]. We compare BLEND with minimap2 and MinHash (i.e., MHAP) by finding overlapping reads. We then generate the assemblies from the overlapping reads to compare the qualities of these assemblies. Second, read mapping uses seeds to find similar portions between a reference genome and a read before performing the read alignment. Aligning a read to a reference genome shows the edit operations (i.e., match, substitution, insertion, and deletions) to make the read identical to the portion of the reference genome, which is useful for downstream analysis (e.g., variant calling [65]). We compare BLEND with minimap2, LRA [66], Winnowmap2, and S-conLSH by mapping long and paired-end short reads to their reference genomes. This paper makes the following **key contributions**:

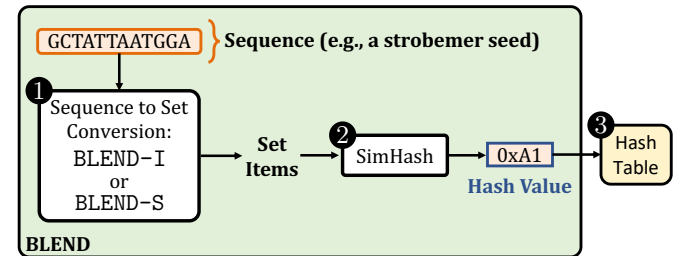
- We introduce BLEND, the *first* mechanism that can quickly and efficiently find *fuzzy* seed matches between sequences with a single lookup.
- We propose two mechanisms for converting seeds into a set of items that the SimHash technique requires: 1) BLEND-I and 2) BLEND-S. We show that BLEND-S provides better speedup and accuracy than BLEND-I when using PacBio HiFi reads for read overlapping and read mapping. When using ONT, PacBio CLR, and short reads, BLEND-I provides significantly better accuracy than BLEND-S with similar performance.
- For read overlapping, we show that BLEND provides speedup compared to minimap2 and MHAP by  $2.6\times$ - $63.5\times$  (on average  $19.5\times$ ),  $22.2\times$ - $6317.1\times$  (on average  $1386.8\times$ ) while reducing the memory overhead by  $0.9\times$ - $9.7\times$  (on average  $3.6\times$ ),  $35.9\times$ - $238.1\times$  (on average  $130.5\times$ ), respectively.

- We show that BLEND usually finds *longer* overlaps between reads while using *fewer* seed matches than other tools, which improves the performance and memory space efficiency for read overlapping.
- We find that we can construct more accurate assemblies with similar contiguity by using the overlapping reads that BLEND finds compared to those that minimap2 finds.
- For read mapping, we show that BLEND provides speedup compared to minimap2, LRA, Winnowmap2, and S-conLSH by  $0.7\times$ - $3.7\times$  (on average  $1.7\times$ ),  $1.7\times$ - $26.5\times$  (on average  $12.5\times$ ),  $1.8\times$ - $19.3\times$  (on average  $7.5\times$ ),  $6.8\times$ - $58.1\times$  (on average  $21.8\times$ ) while maintaining a similar memory overhead by  $0.5\times$ - $1.2\times$  (on average  $0.9\times$ ),  $0.4\times$ - $1.2\times$  (on average  $0.8\times$ ),  $0.9\times$ - $2.7\times$  (on average  $1.7\times$ ),  $0.5\times$ - $5.6\times$  (on average  $1.9\times$ ), respectively.
- We show that BLEND provides a similar read mapping accuracy to minimap2, and Winnowmap2 usually provides the best read mapping accuracy.
- We open source our BLEND implementation as integrated into minimap2.
- We provide the open source SIMD implementation of the SimHash technique that BLEND employs.

## 2. Methods

We propose **BLEND**, a mechanism that can efficiently find fuzzy (i.e., approximate) seed matches with a single lookup of their hash values. To find fuzzy seed matches, BLEND introduces a new mechanism that enables generating the same hash values for highly similar seeds. By combining this mechanism with any seeding approach (e.g., minimizer k-mers or strobemers), BLEND can find fuzzy seed matches between sequences with a single lookup of hash values.

Figure 4 shows the overview of steps to find fuzzy seed matches with a single lookup in three steps. First, BLEND starts with converting the input sequence it receives from a seeding technique (e.g., a strobemer sequence in Figure 3) to its set representation as the SimHash technique generates the hash value of the set using its items ①. To enable effective and efficient integration of seeds with the SimHash technique, BLEND proposes two mechanisms for identifying the items of the set of the input sequence: 1) BLEND-I and 2) BLEND-S. Second, after identifying the items of the set, BLEND uses this set with the SimHash technique to generate the hash value for the input sequence ②. BLEND uses the SimHash technique as it allows for generating the same hash value for highly similar sets. Third, BLEND uses the hash tables as shown in Figure 1 with the hash values it generates to enable finding fuzzy seed matches with a single lookup of their hash values ③.



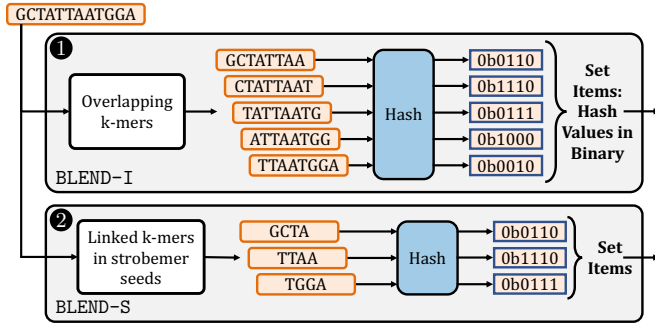
**Figure 4: Overview of BLEND.** ① BLEND uses BLEND-I or BLEND-S for converting a sequence into its set of items. ② BLEND generates the hash value of the input sequence using its set of items with the SimHash technique. ③ BLEND uses hash tables for finding fuzzy seed matches with a single lookup of the hash values that BLEND generates.



## 2.1. Sequence to Set Conversion

Our goal is to convert the input sequences that BLEND receives from any seeding technique (Figure 3) to their proper set representations so that BLEND can use the items of sets for generating the hash values of input sequences with the SimHash technique. To achieve effective and efficient conversion of sequences into their set representations in different scenarios, BLEND provides two mechanisms: 1) BLEND-I and 2) BLEND-S, as we show in Figure 5.

The goal of the first mechanism, BLEND-I, is to provide high sensitivity for a single character change in the input sequences that seeding mechanisms provide when generating their hash values such that two sequences are likely to have the same hash value if they differ by a few characters. BLEND increases the chance of detecting single character changes in input sequences by using *all* the overlapping k-mers of an input sequence for determining the items of its set, as we show in ① of Figure 5. BLEND hashes these k-mers using any hash function and uses the hash values of the k-mers as the items of the set that represents the input sequence. Although BLEND-I can be integrated with any seeding mechanism, we integrate it with the minimizer seeding mechanism, as we show in Figure 3 as proof of work.



**Figure 5: Overview of two mechanisms used for determining the set items of input sequences. ① BLEND-I uses the hash values of all the overlapping k-mers of an input sequence as the set items. ② BLEND-S uses the hash values of only the k-mers selected by the strobemer seeding mechanism.**


The goal of the second mechanism, BLEND-S, is to allow indels and substitutions when matching the input sequences that BLEND receives such that two sequences are likely to have the same hash value if seeds differ by a few k-mers. To achieve this, BLEND-S uses only the selected k-mers that the strobemer-like seeding mechanisms find and link [59], as shown in ② of Figure 5. Since BLEND-S uses only the selected k-mers of strobemer seeds, BLEND enables k-mers to mismatch between input sequences rather than enabling a few characters to mismatch that BLEND-I provides. To ensure the correctness of strobemer seeds when matching them based on their hash values, BLEND-S uses *only* the selected k-mers from the same strand as these k-mers that strobemer-like approaches find may come from either strand strands of DNA. BLEND-S uses the hash values of all such selected k-mers as the set items of the input sequence.

## 2.2. Integrating the SimHash Technique

Our goal is to enable efficient comparisons of equivalence or high similarity between seeds with a single lookup by generating the same hash value for highly similar or equivalent seeds. To enable generating the same hash value for these seeds, BLEND uses the SimHash technique [61]. The SimHash technique takes a set of items and generates a hash value for the set

using its items. The key benefit of the SimHash technique is that it allows generating the same hash value for highly similar sets while enabling any *arbitrary* items to mismatch between sets. To exploit the key benefit of the SimHash technique, BLEND efficiently and effectively integrates the SimHash technique with the set items that BLEND-I or BLEND-S determine. BLEND uses these set items for generating the hash values of seeds such that highly similar seeds can have the same hash value to enable finding fuzzy seed matches with a single lookup of their hash values.

BLEND employs the SimHash technique in three steps: 1) encoding the set items in a vector, 2) performing vector additions, and 3) decoding the vector to generate the hash value for the set that BLEND-I or BLEND-S determine, as we show in Figure 6. To enable efficient computations between vectors, BLEND uses SIMD operations when performing all these three steps.



**Figure 6: The overview of the steps in the SimHash technique for calculating the hash value of a given set of items. The set items are the hash values represented in their binary form. Binary to Vector Encoding converts these set items to their corresponding vector representations. Sum performs the vector additions and stores the result in a separate vector that we call the counter vector. Decoding generates the hash value of the set based on the values in the counter vector. BLEND uses SIMD operations for these three steps, as indicated by SIMD. We highlight in red how 0 bits are converted and propagated in the SimHash technique.**

First, the goal of the *binary to vector encoding* step is to transform hash values of set items into vectors so that BLEND can efficiently perform the bitwise arithmetic operations that the SimHash technique uses. SimHash uses bitwise operations to calculate the difference between the number of 1 and 0 bits at each bit position of hash values, which we call *consensus* calculation. Vector encoding enables performing simple bitwise increments and decrements at each position to calculate the consensus. To perform these increments and decrements efficiently, the vector encoding step generates a vector for each set item where 1 and 0 bits of a hash value are represented as 1 and -1 in the vector, respectively.

Second, the goal of the vector addition operation is to determine the bit positions where the number of 1 bits is greater than the number of 0 bits among the set items, which we call determining the *majority* bits. The key insight in determining these majority bits is that highly similar sets are likely to result in *similar* majority results because a few differences between two similar sets are unlikely to change the majority bits at each position, given that there is a sufficiently large number of items involved in this majority calculation. To efficiently determine the majority of bits at each position, BLEND counts the number of 1 and 0 bits at a position by using the vectors it generates in the vector encoding step, as shown with the addition step (Sum) in Figure 6. The vector addition performs simple additions of +1 or -1 values between the vector elements and stores the result in a separate *counter* vector. The values in this counter vector show the majority of bits at each position of the set items. Since BLEND assigns -1 for 0 bits and 1 for 1 bits, the majority of bits at a position is either 1) 1 if the corresponding value in the counter vector is greater than 0 or 2) otherwise if the values are less than or equal to 0.

Third, to generate the hash value of a set, BLEND uses the majority of bits that it determines by calculating the counter vector. To this end, BLEND decodes the counter vector into a hash value in its binary form, as shown in Figure 6 with the decoding step. The decoding operation is a simple conditional operation where each bit of the final hash value is determined based on its corresponding value at the same position in the counter vector. BLEND assigns the bit either 1) 1 if the value at the corresponding position of the counter vector is greater than 0 or 2) 0 if otherwise. Thus, each bit of the final hash value of the set shows the majority voting result of set items of a seed. We use this final hash value for the input sequence that the seeding techniques provide because highly similar sequences are likely to have many characters or k-mers in common, which essentially leads to generating *similar* set items by using BLEND-I or BLEND-S. Properly identifying the set items of similar sequences enables BLEND to find similar majority voting results with the SimHash technique, which can lead to generating the same final hash value for similar sequences. This enables BLEND to find fuzzy seed matches with a single lookup using these hash values. We explain the details of the reasoning why two similar or same seeds may have the same hash value in Supplementary Section S1 and provide a real example of generating the hash values for two different seeds using different parameter settings in Supplementary Section S1.1 and Supplementary Tables S1- S8.

### 2.3. Using the Hash Tables

Our goal is to enable an efficient lookup of the hash values of seeds to find fuzzy seed matches with a single lookup. To this end, BLEND uses hash tables in two steps. First, BLEND stores the hash values of all the seeds of target sequences (e.g., a reference genome) in a hash table, usually known as the *indexing* step. Keys of the hash table are hash values of seeds and the value that a key returns is a *list* of metadata information (i.e., seed length, position in the target sequence, and the unique name of the target sequence). BLEND keeps minimal metadata information for each seed sufficient to locate seeds in target sequences. Since similar or equivalent seeds can share the same hash value, BLEND stores these seeds using the same hash value in the hash table. Thus, a query to the hash table returns all fuzzy seed matches with the same hash value.

Second, BLEND iterates over all query sequences (e.g., reads) and uses the hash table from the indexing step to find fuzzy seed matches between query and target sequences. The query to the hash table returns the list of seeds of the target sequences that have the same hash value as the seed of a query sequence. Thus, the list of seeds that the hash table returns is the list of fuzzy seed matches for a seed of a query sequence as they share the same hash value. BLEND can find fuzzy seed matches with a single lookup using the hash values it generates for the seeds from both query and target sequences.

BLEND finds fuzzy seed matches mainly for two important genomics applications: read overlapping and read mapping. For these applications, BLEND stores all the list of fuzzy seed matches between query and target sequences to perform *chaining* among fuzzy seed matches that fall in the same target sequence (overlapping reads) optionally followed by alignment (read mapping) as described in minimap2 [15].

## 3. Evaluation

### 3.1. Evaluation Methodology

We replace the mechanism in minimap2 that generates hash values for seeds with BLEND to find fuzzy seed matches when

performing end-to-end read overlapping and read mapping. We also incorporate the BLEND-I and BLEND-S mechanisms in the implementation and provide the user to choose either of these mechanisms when using BLEND. We provide a set of default parameters we optimize based on sequencing technology and the application to perform (e.g., read overlapping). We explain the details regarding the parameters in Supplementary Tables S17 and S18. We determine these default parameters empirically by testing the performance and accuracy of BLEND with different values for some parameters (i.e., k-mer length, number of k-mers to include in a seed, and the window length) as shown in Supplementary Table S12. We show the trade-offs between the seeding mechanisms BLEND-I and BLEND-S in Supplementary Figures S1 and S2 and Supplementary Tables S9 - S11 regarding their performance and accuracy.

For our evaluation, we use real and simulated read datasets as well as their corresponding reference genomes. We list the details of these datasets in Table 1. To evaluate BLEND in several common scenarios in read overlapping and read mapping, we classify our datasets into three categories: 1) highly accurate long reads (i.e., PacBio HiFi), 2) erroneous long reads (i.e., PacBio CLR and Oxford Nanopore Technologies), and 3) short reads (i.e., Illumina). We use PBSIM2 [67] to simulate erroneous reads from both PacBio CLR and Oxford Nanopore Technologies (ONT). HiFi and Illumina reads are from real datasets. To use realistic depth of coverage, we use SeqKit [68] to down-sample the original *E. coli*, and *D. ananassae* reads to 100 $\times$  and 50 $\times$  sequencing depth of coverage, respectively. For *D. ananassae* and *E. coli* genomes, the reference genomes are the high-quality assemblies generated using the same read sets we use for these genomes [69].

**Table 1: Details of datasets used in evaluation.**

Organism	Library	Reads (#)	Seq. Depth	SRA Accession	Reference Genome
Human CHM13	PacBio HiFi	3,167,477	16	SRR11292122-3	GCA_009914755.3
<i>D. ananassae</i>	PacBio HiFi	1,195,370	50	SRR11442117	[69]
<i>E. coli</i>	PacBio HiFi	38,703	100	SRR11434954	[69]
Yeast	PacBio CLR*	270,849	200	Simulated P6-C4	GCA_000146045.2
	Oxford Nanopore Tech.*	135,296	100	Simulated R9.5	GCA_000146045.2
	Illumina MiSeq	3,318,467	80	ERR1938683	

\* We use PBSIM2 to generate the simulated PacBio and ONT reads from the Yeast genome. We include the simulated chemistry under SRA Accession.

We evaluate BLEND based on two use cases: 1) read overlapping and 2) read mapping to a reference genome. For read overlapping, we perform *all-vs-all overlapping* to find all pairs of overlapping reads within the same dataset (i.e., the target and query sequences are the same set of sequences). To evaluate the efficiency of overlaps, we calculate the average length of overlaps and the number of seed matches per overlap to calculate the overlap statistics. To evaluate the quality of overlapping reads based on the accuracy of the assemblies we generate from overlaps, we use miniasm [14]. We use miniasm because it does not perform error correction when generating *de novo* assemblies, which allows us to directly assess the quality of overlaps without using additional approaches for improving the accuracy of assemblies. We use mhaf2paf.pl package as provided by miniasm to convert the output of MHAF to the format miniasm requires (i.e., PAF). We use QUAST [70] to measure statistics related to the contiguity, length, and the accuracy of *de novo* assemblies, such as the overall assembly length, largest contig, NG50 and NGA50 statistics (i.e., statistics related to the length of the shortest contig at the half of the overall reference genome length), k-mer completeness (i.e., amount of shared k-mers between the reference genome and an assembly), the ratio of the misassembly length to the actual assembly length, and GC content (i.e., the ratio of G and C bases to A and T bases in

an assembly). We use dnadiff [71] to measure the accuracy of *de novo* assemblies based on 1) the average identity of an assembly when compared to its reference genome and 2) the fraction of overall bases in a reference genome that align to a given assembly (i.e., genome fraction). We compare BLEND with minimap2 [15] and MHAP [45] for read overlapping. For the Human CHM13 genome, MHAP generates a large output such that we cannot generate the assembly as miniasm exceeds the memory space we have in our system (i.e., 1TB).


For read mapping, we map all reads in a dataset (i.e., query sequences) to their corresponding reference genome (i.e., target sequence). To evaluate the coverage of read mapping, we use BEDTools [72] and Mosdepth [73] to calculate the breadth of coverage (i.e., percentage of bases in a reference genome covered by at least one read) and mean depth of coverage (i.e., the average number of read alignments per base in a reference genome), respectively. To evaluate the quality of read mapping, we use BAMUtil [74] to calculate the mapping rate (i.e., number of aligned reads) and rate of properly paired reads for paired-end mapping. To evaluate the accuracy of read mapping, we measure 1) the overall number of true mappings with very high mapping quality (i.e., 60 mapping quality score) and 2) the overall error rate in read mapping. To generate these results, we use the `paftools` `mapeval` tool provided by minimap2 that takes the reads from PBSIM2 such that the read IDs are annotated with their true mapping information and finds incorrect mappings in the SAM files provided by read mapping tools. We compare BLEND with minimap2, LRA [66], Winnowmap2 [46, 47], and S-conLSH [54, 55] when performing read mapping. We do not evaluate 1) LRA, Winnowmap2, and S-conLSH when mapping paired-end short reads as these tools do not support mapping paired-end short reads, and 2) S-conLSH for the *D. ananassae* genome as S-conLSH crashes due to a segmentation fault when mapping reads to the *D. ananassae* reference genome.

For both use cases, we use the `time` command in Linux to evaluate the performance and peak memory footprints. When applicable, we use the default parameters of all the tools suggested for certain use cases and sequencing technologies (e.g., mapping HiFi reads in minimap2). Since minimap2 and MHAP do not provide default parameters for read overlapping using HiFi reads, we use the parameters that HiCanu [75] uses for overlapping HiFi reads with minimap2 and MHAP. We provide the details regarding the parameters and versions we use for each tool in Supplementary Tables S17, S18, and S19. To show that BLEND can provide better accuracy with the same set of parameters, we use the same parameters that BLEND uses in minimap2 and show the performance and accuracy results in Supplementary Figures S3 and S4 and Supplementary Tables S13 - S15.

### 3.2. Use Case 1: Read Overlapping

**3.2.1. Performance.** Figure 7 shows the CPU time and peak memory footprint comparisons for read overlapping. We make the following five observations. First, BLEND provides speedup by  $2.6\times$ - $63.5\times$  (on average  $19.5\times$ ) and  $22.2\times$ - $6317.1\times$  (on average  $1386.8\times$ ) while reducing the memory footprint by  $0.9\times$ - $9.7\times$  (on average  $3.6\times$ ) and  $35.9\times$ - $238.1\times$  (on average  $130.5\times$ ) compared to minimap2 and MHAP, respectively. BLEND is significantly more performant and provides less memory overheads than MHAP because MHAP generates many hash values for seeds regardless of the length of the sequences, while BLEND provides windowing guarantees when generating seeds, which allows sampling the number of seeds based on the sequence length. Second, when considering only HiFi reads,

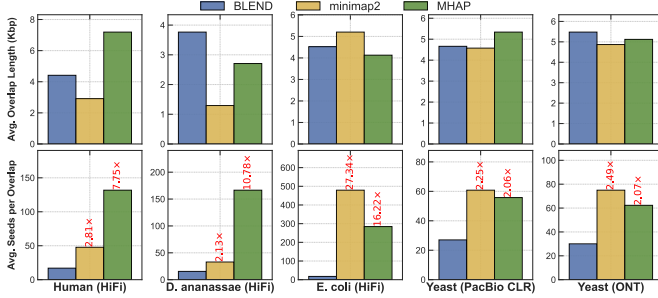
BLEND provides even higher speedup, on average, by  $30.8\times$  and  $2294.4\times$  while reducing the memory footprint by  $5.4\times$  and  $190.4\times$  compared to minimap2 and MHAP, respectively. HiFi reads allow BLEND to increase the window length (i.e.,  $w = 200$ ) when finding the minimizer k-mer of a seed, which improves the performance and reduces the memory overhead without reducing the accuracy. This is possible mainly because BLEND can find *both* fuzzy and exact seed matches, which enables BLEND to find *unique* fuzzy seed matches that minimap2 *cannot* find due to its exact-matching seed requirement. Third, we find that BLEND requires less than 16GB of memory space for *all* the datasets, making it largely possible to find overlapping reads even with a personal computer with relatively small memory space. BLEND has a lower memory footprint because 1) BLEND uses as many seeds as the number of minimizer k-mers per sequence to benefit from the reduced storage requirements that minimizer k-mers provide, and 2) the window length is larger than minimap2 as BLEND can tolerate increasing this window length with the fuzzy seed matches without reducing the accuracy. Fourth, when using erroneous reads (i.e., PacBio CLR and ONT), BLEND provides better performance but similar memory overhead to minimap2. The set of parameters we use for erroneous reads prevents BLEND from using large windows (i.e.,  $w = 10$  instead of  $w = 200$ ) without reducing the accuracy of read overlapping. This causes BLEND to use many seeds, which requires higher memory than when using erroneous reads. Fifth, we use the same set of parameters (i.e., the seed length and the window length) with minimap2 that BLEND uses to observe the benefits that BLEND provides with PacBio CLR and ONT datasets. We cannot perform the same experiment for the HiFi datasets because BLEND uses seeds of length 31, which minimap2 cannot support due to the maximum seed length limitation in its implementation (i.e., max. 28). We call this version of minimap2, minimap2-Eq. We show in Supplementary Figure S3 that minimap2-Eq performs, on average,  $\sim 10\%$  better than BLEND when using the same set of parameters while providing worse accuracy than BLEND when generating the assemblies, as shown in Supplementary Table S13. Performing worse than minimap2-Eq is expected because we apply the same sampling rate and k-mer size and use a much simpler hashing mechanism than BLEND, which leads to using a similar number of seeds while processing them faster with cheaper hashing. The main benefit of BLEND is to provide overall higher accuracy than both the baseline minimap2 and minimap2-Eq, which we can achieve by finding unique fuzzy seed matches that minimap2 cannot find. We conclude that BLEND is significantly more memory-efficient and faster than other tools to find overlaps, especially when using HiFi reads with its ability to sample many seeds using high values of  $w$  without reducing the accuracy.



**Figure 7: CPU time and peak memory footprint comparisons of read overlapping.**



**3.2.2. Overlap Statistics.** Figure 8 shows the average length of overlaps and the average number of seed matches that each tool finds to identify the overlaps between reads. We make the following two key observations. First, we observe that BLEND finds overlaps longer than minimap2 and MHAP can find in most cases. BLEND can 1) uniquely find the fuzzy seed matches that the exact-matching-based tools cannot find and 2) perform chaining on these fuzzy seed matches to increase the length of overlap using many fuzzy seed matches that are relatively close to each other. Finding more distinct seeds and chaining these seeds enable BLEND to find longer overlaps than other tools. Second, BLEND uses significantly fewer seed matches than other tools, up to  $27.34\times$ , to find these longer overlaps. This is mainly because BLEND needs much fewer seeds as it uses 1) higher window lengths than minimap2 and 2) provides windowing guarantees, unlike MHAP. We conclude that the performance and memory-efficiency improvements in read overlapping are proportional to the reduction in the seed matches that BLEND uses to find overlapping reads. Thus, reducing the number of seed matches helps improve the performance and memory space efficiency of BLEND.



**Figure 8: Average length of overlaps and average number of seeds used to find a single overlap.**

**3.2.3. Assembly Quality Assessment.** Our goal is to assess the quality of assemblies generated using the overlapping reads found by BLEND, minimap2, and MHAP. Table 2 shows the statistics related to the accuracy of assemblies (i.e., the 7 statistics on the left-most part of the table) and the statistics related to assembly length and contiguity (i.e., the 3 statistics on the right-most part of the table) when compared to their respective reference genomes. We make the following five key observations based on the accuracy results of assemblies. First, we observe that we can construct more accurate assemblies in terms of average identity and k-mer completeness when we use the overlapping reads that BLEND finds than those minimap2 and MHAP find. These results show that the assemblies we generate using the BLEND overlaps are more similar to their corresponding reference genome. BLEND can find unique and accurate overlaps using fuzzy seed matches that lead to more accurate *de novo* assemblies than the minimap2 and MHAP overlaps due to their lack of support for fuzzy seed matching. Second, we observe that assemblies generated using BLEND overlaps usually cover a higher fraction of the reference genome than minimap2 and MHAP overlaps. Third, although the average identity and genome fraction results seem mixed for the PacBio CLR and ONT reads such that BLEND is best in terms of either average identity or genome fraction, we believe these two statistics should be considered together (e.g., by multiplying both results). This is because a highly accurate but much smaller fraction of the assembly can align to a reference genome, giving the best results for the average identity. We observe that this is the case for the *D. ananassae* and *Yeast* (PacBio CLR) genomes such that MHAP provides a very high average identity only for

the much smaller fraction of the assemblies than the assemblies generated using BLEND and minimap2 overlaps. Thus, when we combine average identity and genome fraction results, we observe that BLEND consistently provides the best results for all the datasets. Fourth, BLEND and minimap2 usually provide the best results when the assemblies are aligned to the reference genome using QUAST in terms of aligned length, misassembly ratio, and NGA50 statistics. In most cases, QUAST cannot generate these statistics for the MHAP results as usually a small portion of the assemblies align the reference genome when the MHAP overlaps are used. Fifth, we find that assemblies generated from BLEND overlaps are less biased than minimap2 and MHAP overlaps. Our observation is based on the average GC content in Table 2 and the GC distributions in Supplementary Figure S5 that are mostly closer to the reference genomes. We conclude that BLEND overlaps yield assemblies with higher accuracy and less bias than the assemblies that the minimap2 and MHAP overlaps generate in most cases.

Table 2 shows the results related to assembly length and contiguity on its right-most part, from which we make the following two observations. First, we show that BLEND yields assemblies with better contiguity when using HiFi reads due to mostly 1) higher NG50 and 2) the largest contigs that BLEND generates are longer than that of minimap2 with the exception of the human genome. BLEND finds unique overlaps that minimap2 cannot find because of BLEND’s ability to detect fuzzy seed matches. We believe such unique overlaps *fill the gap* that minimap2 cannot, which improves the contiguity in most cases. Second, we observe that BLEND and minimap2 result in assemblies where the overall length is mostly closer to the reference genome assembly. We conclude that BLEND enables generating assemblies with contiguity to minimap2. BLEND achieves both high accuracy and high contiguity by using 1) fuzzy seed matches and 2) fewer seeds to find overlapping reads. Fewer overlaps with fuzzy seed matches are likely to provide the assembly graph with fewer edges to remove in transitive reduction and bubbles to collapse in the genome assembly step, which results in more effective overlap construction when using miniasm to generate *de novo* assemblies [14].

### 3.3. Use Case 2: Read Mapping

**3.3.1. Performance.** Figure 9 shows the CPU time and the peak memory footprint comparisons when performing read mapping to the corresponding reference genomes. We make the following two key observations. First, we observe that BLEND provides speedup by  $0.7\times$ - $3.7\times$  (on average  $1.7\times$ ),  $1.7\times$ - $26.5\times$  (on average  $12.5\times$ ),  $1.8\times$ - $19.3\times$  (on average  $7.5\times$ ), and  $6.8\times$ - $58.1\times$  (on average  $21.8\times$ ) compared to minimap2, LRA, Winnowmap2, and S-conLSH, respectively. Although BLEND performs better than most of these tools, the speedup ratio is usually lower than what we observe in read overlapping. Read mapping includes an additional computationally costly step that read overlapping skips, which is the read alignment. The extra overhead of read alignment slightly hinders the benefits that BLEND provides that we observe in read overlapping. Second, we find that BLEND provides slightly more memory overhead than minimap2 and LRA (on average  $0.9\times$  and  $0.8\times$ ) and lower than Winnowmap2 and S-conLSH by  $0.9\times$ - $2.7\times$  (on average  $1.7\times$ ),  $0.5\times$ - $5.6\times$  (on average  $1.9\times$ ), respectively. BLEND cannot provide the similar reductions in the memory overhead that we observe in read overlapping due to the more narrow window length ( $w = 50$  instead of  $w = 200$ ) it uses to find the minimizer k-mers for HiFi reads with high accuracy. Using a narrow window length generates more seeds to store in a hash table than using a high window length, which proportionally

Table 2: Assembly quality comparisons.

Dataset	Tool	Average Identity (%)	Genome Fraction (%)	K-mer Compl. (%)	Aligned Length (Mbp)	Misassembly Ratio (%)	NGA50 (Kbp)	Average GC (%)	Assembly Length (Mbp)	Largest Contig (Mbp)	NG50 (Kbp)
Human CHM13	BLEND	<b>99.8526</b>	<b>98.4847</b>	<b>90.15</b>	3,092.59	<b>0.0108</b>	5,442	<b>40.78</b>	<b>3,095,210</b>	22.840	5,442
	minimap2	99.7421	97.1493	83.05	<b>3,095.49</b>	0.0503	<b>7,133</b>	40.71	3,100,974	<b>47.139</b>	<b>7,134</b>
	MHAP	N/A	N/A	N/A	N/A	N/A	N/A	N/A	N/A	N/A	N/A
	Reference	100	100	100	3,054.83	0	154,260	40.85	3,054,832	248.387	154,260
<i>D. ananassae</i>	BLEND	<b>99.7856</b>	<b>97.2308</b>	<b>86.43</b>	240.39	0.1230	<b>792</b>	<b>41.75</b>	<b>247,153</b>	<b>6.233</b>	<b>799</b>
	minimap2	99.7044	96.3190	72.33	<b>289.45</b>	<b>0.1013</b>	273	41.68	298,280	4.434	273
	MHAP	99.5551	0.7276	0.21	2.29	0.2197	N/A	42.07	2,350	0.286	N/A
	Reference	100	100	100	213.81	0	26,427	41.81	213,818	30.673	26,427
<i>E. coli</i>	BLEND	<b>99.8320</b>	<b>99.8801</b>	<b>87.91</b>	<b>5.12</b>	<b>0.0344</b>	<b>3,417</b>	<b>50.53</b>	5.122	<b>3.417</b>	<b>3,417</b>
	minimap2	99.7064	99.8748	79.27	5.09	0.3068	3,087	50.47	<b>5.042</b>	3.089	3,089
	MHAP	N/A	N/A	N/A	N/A	N/A	N/A	N/A	5.094	N/A	N/A
	Reference	100	100	100	5.05	0	4,945	50.52	5.046	4.946	4,945
Yeast (PacBio)	BLEND	89.1582	<b>97.6297</b>	<b>11.13</b>	<b>0.227</b>	N/A	N/A	<b>38.80</b>	13.679	1.105	551
	minimap2	88.9002	96.9709	9.74	0.195	N/A	N/A	38.85	<b>12.333</b>	<b>1.561</b>	<b>828</b>
	MHAP	<b>89.2182</b>	88.5928	9.5	0.186	N/A	N/A	38.81	10.990	1.024	436
	Reference	100	100	100	12.16	0	924	38.15	12.157	1.532	924
Yeast (ONT)	BLEND	<b>89.7622</b>	99.2982	<b>13.68</b>	<b>0.377</b>	N/A	N/A	<b>38.66</b>	<b>12.164</b>	1.554	825
	minimap2	88.9393	<b>99.6878</b>	12.06	0.328	N/A	N/A	38.74	12.373	<b>1.560</b>	<b>942</b>
	MHAP	89.1970	89.2785	11.35	0.297	N/A	N/A	38.84	10.920	1.443	619
	Reference	100	100	100	12.16	0	924	38.15	12.157	1.532	924

Best results are highlighted with **bold** text. For most metrics, the best results are the ones closest to the corresponding value of the reference genome.

The best results for *Aligned Length* are determined by the highest number within each dataset. We do not highlight the reference results as the best results.

N/A indicates that we could not generate the corresponding result because either the tool failed or QUAST failed to generate the statistic.

increases the peak memory space requirements. Third, we use the same parameters that BLEND uses with minimap2 to assess the accuracy benefits of finding fuzzy seed matches rather than just using exact-matching seeds in read mapping. We call the version of minimap2 that uses the same set of parameters minimap2-Eq. We observe that BLEND performs, on average,  $1.3\times$  better than minimap2-Eq, as shown in Supplementary Figure S4, while mostly providing higher read mapping quality and accuracy, as shown in Supplementary Tables S14 and S15. We conclude that BLEND, on average, performs better than all tools without reducing the accuracy and provides a better memory footprint than Winnowmap2 and S-conLSH.




Figure 9: CPU time and peak memory footprint comparisons of read mapping.

**3.3.2. Read Mapping Accuracy.** To assess the accuracy of read mapping, we map two simulated erroneous read sets (i.e., PacBio CLR and ONT) to their reference genomes and compare the read mapping results with their true mapping results as generated by PBSIM2. To this end, we use paftools to assess the read mapping accuracy and show these results in Supplementary Table S16. We make two observations. First, we observe that BLEND provides the most high-quality true mappings for the PacBio CLR reads, and the lowest overall error rate for the ONT reads. For the cases where BLEND does not provide the best results, minimap2 provides the most high-quality true mappings for the ONT reads, and Winnowmap2 provides the best error rate for the PacBio CLR reads. Second, LRA and S-conLSH map most of the reads to incorrect locations, making these tools

inaccurate for mapping erroneous reads compared to BLEND, minimap2, and Winnowmap2. We conclude that although the results are mixed, BLEND is the only tool that always performs best in either of these two metrics, providing the overall best accuracy results as generated by the paftools.

**3.3.3. Read Mapping Quality.** Our goal is to assess the quality of read mappings in terms of four metrics: mean depth of coverage, breadth of coverage, number of aligned reads, and the ratio of the paired-end reads that are properly paired in mapping. Table 3 shows the quality of read mappings based on these metrics when using BLEND, minimap2, LRA, and Winnowmap2. We exclude S-conLSH from the read mapping quality comparisons as we cannot convert its SAM output to BAM format to properly index the BAM file due to the issues with its SAM output format. We make three observations. First, BLEND, minimap2, and Winnowmap2 cover most portion of the reference genomes after read mapping than LRA does. This result shows that these tools are less biased in mapping reads to particular regions (e.g., repetitive regions) than LRA. Second, both BLEND and minimap2 map an almost complete set of reads to the reference genome for all the datasets, while Winnowmap2 suffers from a slightly lower number of aligned reads when mapping erroneous PacBio CLR and ONT reads. This result shows that although Winnowmap2 generates lower error rates for the PacBio CLR reads, the higher portion of the unmapped reads is not counted towards the error rate. This suggests us Winnowmap2 provides a better precision while BLEND and minimap2 provide a better recall in terms of the read mapping accuracy. Third, we find that all the tools generate read mappings with a depth of coverage significantly close to their sequencing depth of coverage. This shows that almost all reads map to the reference genome evenly. We conclude that read mapping qualities of BLEND, minimap2, and Winnowmap2 are highly similar, while LRA provides slightly worse results. It is worth noting that BLEND can achieve comparable accuracy while using larger window lengths when finding the minimizer k-mers, which usually has an effect on the accuracy. BLEND does this by finding unique fuzzy seed matches that the other tools cannot find due to their exact-matching seed requirements.



**Table 3: Read mapping quality comparisons.**

Dataset	Tool	Mean Depth of Cov. ( $\times$ )	Breadth of Coverage (%)	Aligned Reads (#)	Properly Paired (%)
<i>Human CHM13</i>	BLEND	<b>16.58</b>	<b>99.99</b>	3,171,916	NA
	minimap2	<b>16.58</b>	<b>99.99</b>	<b>3,172,261</b>	NA
	LRA	16.37	99.06	3,137,631	NA
	Winnowmap2	<b>16.58</b>	<b>99.99</b>	3,171,313	NA
<i>D. ananassae</i>	BLEND	57.37	99.66	1,223,388	NA
	minimap2	<b>57.57</b>	<b>99.67</b>	1,245,931	NA
	LRA	57.06	99.60	1,235,098	NA
	Winnowmap2	57.40	99.66	<b>1,249,575</b>	NA
<i>E. coli</i>	BLEND	<b>99.14</b>	99.90	39,048	NA
	minimap2	<b>99.14</b>	99.90	<b>39,065</b>	NA
	LRA	99.10	99.90	39,063	NA
	Winnowmap2	<b>99.14</b>	99.90	39,036	NA
<i>Yeast (PacBio)</i>	BLEND	195.84	<b>99.98</b>	269,804	NA
	minimap2	<b>195.86</b>	<b>99.98</b>	<b>269,935</b>	NA
	LRA	194.65	99.97	267,399	NA
	Winnowmap2	192.35	<b>99.98</b>	259,073	NA
<i>Yeast (ONT)</i>	BLEND	97.86	<b>99.97</b>	134,721	NA
	minimap2	<b>97.88</b>	99.96	<b>134,885</b>	NA
	LRA	97.25	99.95	132,862	NA
	Winnowmap2	97.04	99.96	130,978	NA
<i>Yeast (Illumina)</i>	BLEND	<b>79.93</b>	99.97	<b>6,494,489</b>	95.89
	minimap2	79.91	99.97	6,492,994	95.89

Best results are highlighted with **bold** text.

Properly paired rate is only available for paired-end Illumina reads.

## 4. Discussion

We demonstrate that there are usually too many redundant *short* and *exact-matching* seeds used to find overlaps between sequences, as shown in Figure 8. These redundant seeds usually exacerbate the performance and peak memory space requirement problems that read overlapping and read mapping suffer from as the number of chaining and alignment operations proportionally increases with the number of seed matches between sequences [15]. Such redundant computations have been one of the main limitations against developing population-scale genomics analysis due to the high runtime of a single high coverage genome analysis.

There has been a clear interest in using long or fuzzy seed matches because of their potential to find similarities between target and query sequences efficiently and accurately [40]. To achieve this, earlier works mainly focus on either 1) chaining the exact k-mer matches by tolerating the gaps between them to increase the seed region or 2) linking multiple consecutive minimizer k-mers such as strobemer seeds. Chaining algorithms are becoming a bottleneck in read mappers as the complexity of chaining is determined by the number of seed matches [76]. Linking multiple minimizer k-mers enables tolerating indels when finding the matches of short subsequences between genomic sequence pairs, but these seeds (e.g., strobemer seeds) should still exactly match due to the nature of the hash functions used to generate the hash values of seeds. This requires the seeding techniques to generate exactly the same seed to find either exact-matching or approximate matches of short subsequences. We state that any arbitrary k-mer in the seeds should be tolerated to mismatch to improve the sensitivity of any seeding technique, which has a potential for finding more matching regions while using fewer seeds. Thus, we believe BLEND solves the main limitation of earlier works such that it can generate the same hash value for highly similar seeds to find fuzzy seed matches with a single lookup while improving the performance, memory overhead, and accuracy of the applications that use seeds.

We hope that BLEND advances the field and prospers future work in several ways, which we list some of them next. First, we observe that BLEND is *most effective* when using high coverage and highly accurate long reads. Thus, we believe that BLEND is already ready to scale for longer and more accurate sequencing reads. Second, we believe the vector operations are suitable for hardware acceleration to improve the performance of BLEND further. Such an acceleration is mainly useful when a massive amount of k-mers in a seed are used to generate the hash value for a seed, as these calculations can be done in parallel. We already provide the SIMD implementation to calculate the hash values BLEND. We encourage implementing our mechanism for the applications that use seeds with suitable architectures such as processing-in-memory [35, 77], GPUs [78], and SSDs [36] to exploit the massive amount of embarrassingly parallel bitwise operations in BLEND. Third, we believe it is possible to apply the hashing technique we use in BLEND for many seeding techniques with a proper design. We already show we can apply SimHash in regular minimizer k-mers or strobemers. Strobemers can be generated using k-mer sampling strategies other than minimizer k-mers, which are based on syncmers and random selection of k-mers (i.e., randstobes) [60]. It is worth exploring and rethinking the hash functions used in these seeding techniques. Fourth, we believe potential machine learning applications can be used to generate more sensitive hash values, similar to learning-to-hash approaches [79] while using the hash values that BLEND generates as a starting point.

## 5. Conclusion

We propose BLEND, a mechanism that can efficiently find fuzzy seed matches between sequences to improve the performance, memory space efficiency, and accuracy of two important applications significantly: 1) read overlapping, and 2) read mapping. Based on the experiments we perform using real and simulated datasets, we make five key observations. First, BLEND provides significant speedup for read overlapping by  $2.6\times$ - $63.5\times$  (on average  $19.5\times$ ) and  $22.2\times$ - $6317.1\times$  (on average  $1386.8\times$ ) while reducing the peak memory footprint by  $0.9\times$ - $9.7\times$  (on average  $3.6\times$ ) and  $35.9\times$ - $238.1\times$  (on average  $130.5\times$ ) compared to minimap2 and MHAP. Second, we observe that BLEND finds longer overlaps, in general, while using significantly fewer seed matches by up to  $27.34\times$  to find these overlaps. Third, we find that we can usually generate more *accurate* assemblies when using the overlaps that BLEND finds than those found by minimap2 and MHAP. Fourth, for read mapping, we find that BLEND, on average, provides speedup by 1)  $1.7\times$ ,  $12.5\times$ ,  $7.5\times$ , and  $21.8\times$  compared to minimap2, LRA, Winnowmap2, and S-conLSH, respectively. Fifth, we observe that BLEND, minimap2, and Winnowmap2 provide both high quality and better accuracy in read mapping in all datasets, while the accuracy of LRA and S-conLSH is low when using erroneous long reads. We conclude that BLEND can use fewer fuzzy seed matches to significantly improve the performance and reduce the memory overhead of read overlapping without losing accuracy, while BLEND, on average, provides better performance and similar memory footprint in read mapping without reducing the read mapping quality and accuracy.

## 6. Acknowledgements

We thank the SAFARI members for feedback. We are grateful for the detailed comments that Kristoffer Sahlin provided, which improved our mechanism and the manuscript greatly.

## 7. Funding

This work is supported by gifts from Intel [to O.M.]; VMware [to O.M.]; and a Semiconductor Research Corporation grant [to O.M.].

## References

- [1] J. Shendure, S. Balasubramanian, G. M. Church, W. Gilbert, J. Rogers, J. A. Schloss, and R. H. Waterston, "DNA sequencing at 40: past, present and future," *Nature*, vol. 550, Oct. 2017.
- [2] M.-M. Aynaud, J. J. Hernandez, S. Barutcu, U. Braunschweig, K. Chan, J. D. Pearson, D. Trcka, S. L. Prosser, J. Kim, M. Barrios-Rodiles, M. Jen, S. Song, J. Shen, C. Bruce, B. Hazlett, S. Poutanen, L. Attisano, R. Bremner, B. J. Blencowe, T. Mazzulli, H. Han, L. Pelletier, and J. L. Wrana, "A multiplexed, next generation sequencing platform for high-throughput detection of SARS-CoV-2," *Nature Communications*, vol. 12, Mar. 2021.
- [3] G. A. Logsdon, M. R. Vollger, and E. E. Eichler, "Long-read human genome sequencing and its applications," *Nature Reviews Genetics*, vol. 21, Oct. 2020.
- [4] T. Mantere, S. Kersten, and A. Hoischen, "Long-Read Sequencing Emerging in Medical Genetics," *Frontiers in Genetics*, vol. 10, 2019.
- [5] J. M. Friedman, Y. Bombard, M. C. Cornel, C. V. Fernandez, A. K. Junker, S. E. Plon, Z. Stark, B. M. Knoppers, and for the Paediatric Task Team of the Global Alliance for Genomics and Health Regulatory and Ethics Work Stream, "Genome-wide sequencing in acutely ill infants: genomic medicine's critical application?" *Genetics in Medicine*, vol. 21, Feb. 2019.
- [6] J. D. Merker, A. M. Wenger, T. Sneddon, M. Grove, Z. Zappala, L. Fresard, D. Waggott, S. Utiramerur, Y. Hou, K. S. Smith, S. B. Montgomery, M. Wheeler, J. G. Buchan, C. C. Lambert, K. S. Eng, L. Hickey, J. Korlach, J. Ford, and E. A. Ashley, "Long-read genome sequencing identifies causal structural variation in a Mendelian disease," *Genetics in Medicine*, vol. 20, Jan. 2018.
- [7] C. Alkan, B. P. Coe, and E. E. Eichler, "Genome structural variation discovery and genotyping," *Nature Reviews Genetics*, vol. 12, May 2011.
- [8] S. Goodwin, J. D. McPherson, and W. R. McCombie, "Coming of age: ten years of next-generation sequencing technologies," *Nature Reviews Genetics*, vol. 17, may 2016.
- [9] N. Stoler and A. Nekrutenko, "Sequencing error profiles of Illumina sequencing instruments," *NAR Genomics and Bioinformatics*, vol. 3, Mar. 2021.
- [10] H. Zhang, C. Jain, and S. Aluru, "A comprehensive evaluation of long read error correction methods," *BMC Genomics*, vol. 21, Dec. 2020.
- [11] T. Hon, K. Mars, G. Young, Y.-C. Tsai, J. W. Karalius, J. M. Landolin, N. Maurer, D. Kudrna, M. A. Hardigan, C. C. Steiner, S. J. Knapp, D. Ware, B. Shapiro, P. Peluso, and D. R. Rank, "Highly accurate long-read HiFi sequencing data for five complex genomes," *Scientific Data*, vol. 7, Nov. 2020.
- [12] X. Ma, Y. Shao, L. Tian, D. A. Flasch, H. L. Mulder, M. N. Edmonson, Y. Liu, X. Chen, S. Newman, J. Nakitandwe, Y. Li, B. Li, S. Shen, Z. Wang, S. Shurtleff, L. L. Robison, S. Levy, J. Easton, and J. Zhang, "Analysis of error profiles in deep next-generation sequencing data," *Genome Biology*, vol. 20, Mar. 2019.
- [13] D. Senol Cali, J. S. Kim, S. Ghose, C. Alkan, and O. Mutlu, "Nanopore sequencing technology and tools for genome assembly: computational analysis of the current state, bottlenecks and future directions," *Briefings in Bioinformatics*, vol. 20, Jul. 2019.
- [14] H. Li, "Minimap and miniasm: fast mapping and de novo assembly for noisy long sequences," *Bioinformatics*, vol. 32, Jul. 2016.
- [15] H. Li, "Minimap2: pairwise alignment for nucleotide sequences," *Bioinformatics*, vol. 34, Sep. 2018.
- [16] S. Canzar and S. L. Salzberg, "Short Read Mapping: An Algorithmic Tour," *Proceedings of the IEEE*, vol. 105, Mar. 2017.
- [17] J. S. Kim, C. Firtina, M. B. Cavlak, D. Senol Cali, N. Hajinazar, M. Alser, C. Alkan, and O. Mutlu, "AirLift: A Fast and Comprehensive Technique for Remapping Alignments between Reference Genomes," *bioRxiv*, Jan. 2021.
- [18] J. S. Kim, C. Firtina, M. B. Cavlak, D. Senol Cali, C. Alkan, and O. Mutlu, "FastRemap: A Tool for Quickly Remapping Reads between Genome Assemblies," *arXiv*, Jan. 2022.
- [19] B. Ekim, B. Berger, and R. Chikhi, "Minimizer-space de Bruijn graphs: Whole-genome assembly of long reads in minutes on a personal computer," *Cell Systems*, vol. 12, Oct. 2021.
- [20] H. Cheng, G. T. Concepcion, X. Feng, H. Zhang, and H. Li, "Haplotype-resolved de novo assembly using phased assembly graphs with hifiasm," *Nature Methods*, vol. 18, Feb. 2021.
- [21] G. Robertson, J. Schein, R. Chiu, R. Corbett, M. Field, S. D. Jackman, K. Mungall, S. Lee, H. M. Okada, J. Q. Qian, M. Griffith, A. Raymond, N. Thiessen, T. Cezard, Y. S. Butterfield, R. Newsome, S. K. Chan, R. She, R. Varhol, B. Kamoh, A.-L. Prabhu, A. Tam, Y. Zhao, R. A. Moore, M. Hirst, M. A. Marra, S. J. M. Jones, P. A. Hoodless, and I. Birol, "De novo assembly and analysis of RNA-seq data," *Nature Methods*, vol. 7, Nov. 2010.
- [22] F. Meyer, A. Fritz, Z.-L. Deng, D. Koslicki, A. Gurevich, G. Robertson, M. Alser, D. Antipov, F. Beghini, D. Bertrand *et al.*, "Critical Assessment of Metagenome Interpretation-the second round of challenges," *bioRxiv*, 2021.
- [23] N. LaPierre, M. Alser, E. Eskin, D. Koslicki, and S. Mangul, "Metalign: efficient alignment-based metagenomic profiling via containment min hash," *Genome biology*, vol. 21, 2020.
- [24] D. E. Wood, J. Lu, and B. Langmead, "Improved metagenomic analysis with Kraken 2," *Genome Biology*, vol. 20, Nov. 2019.
- [25] C. Firtina, J. S. Kim, M. Alser, D. Senol Cali, A. E. Cicek, C. Alkan, and O. Mutlu, "Apollo: a sequencing-technology-independent, scalable and accurate assembly polishing algorithm," *Bioinformatics*, vol. 36, Jun. 2020.
- [26] R. Vaser, I. Sović, N. Nagarajan, and M. Šikić, "Fast and accurate de novo genome assembly from long uncorrected reads," *Genome Research*, vol. 27, May 2017.
- [27] N. J. Loman, J. Quick, and J. T. Simpson, "A complete bacterial genome assembled de novo using only nanopore sequencing data," *Nature Methods*, vol. 12, Aug. 2015.
- [28] M. Alser, Z. Bingöl, D. S. Cali, J. S. Kim, S. Ghose, C. Alkan, and O. Mutlu, "Accelerating Genome Analysis: A Primer on an Ongoing Journey," *IEEE Micro*, vol. 40, Oct. 2020.
- [29] G. Singh, M. Alser, D. Senol Cali, D. Diamantopoulos, J. Gómez-Luna, H. Corporaal, and O. Mutlu, "FPGA-Based Near-Memory Acceleration of Modern Data-Intensive Applications," *IEEE Micro*, vol. 41, Aug. 2021.
- [30] M. Alser, T. Shahroodi, J. Gómez-Luna, C. Alkan, and O. Mutlu, "SneakySnake: a fast and accurate universal genome pre-alignment filter for CPUs, GPUs and FPGAs," *Bioinformatics*, vol. 36, Dec. 2020.
- [31] S. Angizi, J. Sun, W. Zhang, and D. Fan, "PIM-Aligner: A Processing-in-MRAM Platform for Biological Sequence Alignment," in *2020 Design, Automation Test in Europe Conference Exhibition (DATE)*, 2020.
- [32] S. D. Goenka, Y. Turakhia, B. Paten, and M. Horowitz, "SegAlign: A Scalable GPU-Based Whole Genome Aligner," in *SC20: International Conference for High Performance Computing, Networking, Storage and Analysis*, 2020.
- [33] D. Senol Cali, G. S. Kalsi, Z. Bingöl, C. Firtina, L. Subramanian, J. S. Kim, R. Ausavarungnirun, M. Alser, J. Gomez-Luna, A. Boroumand, A. Norion, A. Scibisz, S. Subramoneyon, C. Alkan, S. Ghose, and O. Mutlu, "GenASM: A High-Performance, Low-Power Approximate String Matching Acceleration Framework for Genome Sequence Analysis," in *2020 53rd Annual IEEE/ACM International Symposium on Microarchitecture (MICRO)*, 2020.
- [34] Y. Turakhia, G. Bejerano, and W. J. Dally, "Darwin: A Genomics Co-Processor Provides up to 15,000X Acceleration on Long Read Assembly," *SIGPLAN Not.*, vol. 53, Mar. 2018.

- [35] J. S. Kim, D. Senol Cali, H. Xin, D. Lee, S. Ghose, M. Alser, H. Hassan, O. Ergin, C. Alkan, and O. Mutlu, "GRIM-Filter: Fast seed location filtering in DNA read mapping using processing-in-memory technologies," *BMC Genomics*, vol. 19, May 2018.
- [36] N. Mansouri Ghiasi, J. Park, H. Mustafa, J. Kim, A. Olgun, A. Gollwitzer, D. Senol Cali, C. Firtina, H. Mao, N. Almadhoun Alserr, R. Ausavarungnirun, N. Vijaykumar, M. Alser, and O. Mutlu, "GenStore: A High-Performance in-Storage Processing System for Genome Sequence Analysis," in *Proceedings of the 27th ACM International Conference on Architectural Support for Programming Languages and Operating Systems*, ser. ASPLOS 2022. New York, NY, USA: Association for Computing Machinery, 2022.
- [37] A. Nag, C. N. Ramachandra, R. Balasubramonian, R. Stutsman, E. Giacomini, H. Kambalasubramanyam, and P.-E. Gaillardon, "GenCache: Leveraging In-Cache Operators for Efficient Sequence Alignment," in *Proceedings of the 52nd Annual IEEE/ACM International Symposium on Microarchitecture*, ser. MICRO '52. New York, NY, USA: Association for Computing Machinery, 2019.
- [38] D. Senol Cali, K. Kanellopoulos, J. Lindegger, Z. Bingöl, G. S. Kalsi, Z. Zuo, C. Firtina, M. B. Cavlak, J. Kim, N. M. Ghiasi, G. Singh, J. Gómez-Luna, N. A. Alserr, M. Alser, S. Subramoney, C. Alkan, S. Ghose, and O. Mutlu, "SeGraM: A Universal Hardware Accelerator for Genomic Sequence-to-Graph and Sequence-to-Sequence Mapping," in *Proceedings of the 49th Annual International Symposium on Computer Architecture*, ser. ISCA '22. New York, NY, USA: Association for Computing Machinery, 2022.
- [39] T. Shahroodi, M. Zahedi, C. Firtina, M. Alser, S. Wong, O. Mutlu, and S. Hamdioui, "Demeter: A fast and energy-efficient food profiler using hyperdimensional computing in memory," *arXiv*, Jun. 2022.
- [40] M. Alser, J. Rotman, D. Deshpande, K. Taraszka, H. Shi, P. I. Baykal, H. T. Yang, V. Xue, S. Knyazev, B. D. Singer, B. Balliu, D. Koslicki, P. Skums, A. Zelikovsky, C. Alkan, O. Mutlu, and S. Mangul, "Technology dictates algorithms: recent developments in read alignment," *Genome Biology*, vol. 22, Aug. 2021.
- [41] M. Alser, J. Lindegger, C. Firtina, N. Almadhoun, H. Mao, G. Singh, J. Gomez-Luna, and O. Mutlu, "Going from molecules to genomic variations to scientific discovery: Intelligent algorithms and architectures for intelligent genome analysis," *arXiv*, May 2022.
- [42] S. Baichoo and C. A. Ouzounis, "Computational complexity of algorithms for sequence comparison, short-read assembly and genome alignment," *Biosystems*, vol. 156-157, Jun. 2017.
- [43] M. Roberts, W. Hayes, B. R. Hunt, S. M. Mount, and J. A. Yorke, "Reducing storage requirements for biological sequence comparison," *Bioinformatics*, vol. 20, Dec. 2004.
- [44] S. Schleimer, D. S. Wilkerson, and A. Aiken, "Winnowing: local algorithms for document fingerprinting," in *Proceedings of the 2003 ACM SIGMOD international conference on Management of data*, 2003.
- [45] K. Berlin, S. Koren, C.-S. Chin, J. P. Drake, J. M. Landolin, and A. M. Phillippy, "Assembling large genomes with single-molecule sequencing and locality-sensitive hashing," *Nature Biotechnology*, vol. 33, Jun. 2015.
- [46] C. Jain, A. Rhie, N. F. Hansen, S. Koren, and A. M. Phillippy, "Long-read mapping to repetitive reference sequences using Winnowmap2," *Nature Methods*, Apr. 2022.
- [47] C. Jain, A. Rhie, H. Zhang, C. Chu, B. P. Walenz, S. Koren, and A. M. Phillippy, "Weighted minimizer sampling improves long read mapping," *Bioinformatics*, vol. 36, Jul. 2020.
- [48] D. DeBlasio, F. Gbosibo, C. Kingsford, and G. Marçais, "Practical Universal K-Mer Sets for Minimizer Schemes," in *Proceedings of the 10th ACM International Conference on Bioinformatics, Computational Biology and Health Informatics*, ser. BCB '19. New York, NY, USA: Association for Computing Machinery, 2019.
- [49] H. Xin, M. Shao, and C. Kingsford, "Context-aware seeds for read mapping," *Algorithms for Molecular Biology*, vol. 15, May 2020.
- [50] H. Xin, D. Lee, F. Hormozdiari, S. Yedkar, O. Mutlu, and C. Alkan, "Accelerating read mapping with FastHASH," *BMC Genomics*, vol. 14, Jan. 2013.
- [51] B. Ma, J. Tromp, and M. Li, "PatternHunter: faster and more sensitive homology search," *Bioinformatics*, vol. 18, Mar. 2002.
- [52] H. Lin, Z. Zhang, M. Q. Zhang, B. Ma, and M. Li, "ZOOM! Zillions of oligos mapped," *Bioinformatics*, vol. 24, Nov. 2008.
- [53] M. David, M. Dzamba, D. Lister, L. Ilie, and M. Brudno, "SHRIMP2: Sensitive yet Practical Short Read Mapping," *Bioinformatics*, vol. 27, Apr. 2011.
- [54] A. Chakraborty and S. Bandyopadhyay, "conLSH: Context based Locality Sensitive Hashing for mapping of noisy SMRT reads," *Computational Biology and Chemistry*, vol. 85, Apr. 2020.
- [55] A. Chakraborty, B. Morgenstern, and S. Bandyopadhyay, "S-conLSH: alignment-free gapped mapping of noisy long reads," *BMC Bioinformatics*, vol. 22, Feb. 2021.
- [56] E. Petrucci, L. Noé, C. Pizzi, and M. Comin, "Iterative Spaced Seed Hashing: Closing the Gap Between Spaced Seed Hashing and k-mer Hashing," *Journal of Computational Biology*, vol. 27, Feb. 2020.
- [57] A. Mallik and L. Ilie, "ALeS: adaptive-length spaced-seed design," *Bioinformatics*, vol. 37, May 2021.
- [58] C.-S. Chin and A. Khalak, "Human Genome Assembly in 100 Minutes," *bioRxiv*, Jan. 2019.
- [59] K. Sahlin, "Effective sequence similarity detection with strobemers," *Genome Research*, vol. 31, Nov. 2021.
- [60] K. Sahlin, "Faster short-read mapping with strobemer seeds constructed from syncmers," *bioRxiv*, Jan. 2021.
- [61] M. S. Charikar, "Similarity Estimation Techniques from Rounding Algorithms," in *Proceedings of the Thirty-Fourth Annual ACM Symposium on Theory of Computing*, ser. STOC '02. New York, NY, USA: Association for Computing Machinery, 2002.
- [62] G. S. Manku, A. Jain, and A. Das Sarma, "Detecting Near-Duplicates for Web Crawling," in *Proceedings of the 16th International Conference on World Wide Web*, ser. WWW '07. New York, NY, USA: Association for Computing Machinery, 2007.
- [63] P. Jaccard, "Nouvelles recherches sur la distribution florale," *Bull. Soc. Vaud. Sci. Nat.*, vol. 44, 1908.
- [64] M. Pop, A. Phillippy, A. L. Delcher, and S. L. Salzberg, "Comparative genome assembly," *Briefings in Bioinformatics*, vol. 5, Sep. 2004.
- [65] A. McKenna, M. Hanna, E. Banks, A. Sivachenko, K. Cibulskis, A. Kernytsky, K. Garimella, D. Altshuler, S. Gabriel, M. Daly, and M. A. DePristo, "The Genome Analysis Toolkit: A MapReduce framework for analyzing next-generation DNA sequencing data," *Genome Research*, vol. 20, Sep. 2010.
- [66] J. Ren and M. J. P. Chaisson, "Ira: A long read aligner for sequences and contigs," *PLOS Computational Biology*, vol. 17, Jun. 2021.
- [67] Y. Ono, K. Asai, and M. Hamada, "PBSIM2: a simulator for long-read sequencers with a novel generative model of quality scores," *Bioinformatics*, vol. 37, Mar. 2021.
- [68] W. Shen, S. Le, Y. Li, and F. Hu, "SeqKit: A Cross-Platform and Ultrafast Toolkit for FASTA/Q File Manipulation," *PLOS ONE*, vol. 11, Oct. 2016.
- [69] E. S. Tvedte, M. Gasser, B. C. Sparklin, J. Michalski, C. E. Hjelman, J. S. Johnston, X. Zhao, R. Bromley, L. J. Tallon, L. Sadzewicz, D. A. Rasko, and J. C. Dunning Hotopp, "Comparison of long-read sequencing technologies in interrogating bacteria and fly genomes," *G3 Genes|Genomes|Genetics*, vol. 11, Jun. 2021.
- [70] A. Gurevich, V. Saveliev, N. Vyahhi, and G. Tesler, "QUAST: quality assessment tool for genome assemblies," *Bioinformatics*, vol. 29, Apr. 2013.



- [71] G. Marçais, A. L. Delcher, A. M. Phillippy, R. Coston, S. L. Salzberg, and A. Zimin, "MUMmer4: A fast and versatile genome alignment system," *PLOS Computational Biology*, vol. 14, Jan. 2018.
- [72] A. R. Quinlan and I. M. Hall, "BEDTools: a flexible suite of utilities for comparing genomic features," *Bioinformatics*, vol. 26, Mar. 2010.
- [73] B. S. Pedersen and A. R. Quinlan, "Mosdepth: quick coverage calculation for genomes and exomes," *Bioinformatics*, vol. 34, Mar. 2018.
- [74] G. Jun, M. K. Wing, G. R. Abecasis, and H. M. Kang, "An efficient and scalable analysis framework for variant extraction and refinement from population scale DNA sequence data," *Genome Research*, Apr. 2015.
- [75] S. Nurk, B. P. Walenz, A. Rhie, M. R. Vollger, G. A. Logsdon, R. Grothe, K. H. Miga, E. E. Eichler, A. M. Phillippy, and S. Koren, "HiCanu: accurate assembly of segmental duplications, satellites, and allelic variants from high-fidelity long reads," *bioRxiv*, Jan. 2020.
- [76] L. Guo, J. Lau, Z. Ruan, P. Wei, and J. Cong, "Hardware acceleration of long read pairwise overlapping in genome sequencing: A race between fpga and gpu," in *2019 IEEE 27th Annual International Symposium on Field-Programmable Custom Computing Machines (FCCM)*, 2019.
- [77] D. Senol Cali, G. S. Kalsi, Z. Bingöl, C. Firtina, L. Subramanian, J. S. Kim, R. Ausavarungnirun, M. Alser, J. Gomez-Luna, A. Boroumand, A. Norion, A. Scibisz, S. Subramoneyon, C. Alkan, S. Ghose, and O. Mutlu, "GenASM: A High-Performance, Low-Power Approximate String Matching Acceleration Framework for Genome Sequence Analysis," in *2020 53rd Annual IEEE/ACM International Symposium on Microarchitecture (MICRO)*, 2020.
- [78] A. Zeni, G. Guidi, M. Ellis, N. Ding, M. D. Santambrogio, S. Hofmeyr, A. Buluç, L. Olikier, and K. Yelick, "LOGAN: High-Performance GPU-Based X-Drop Long-Read Alignment," in *2020 IEEE International Parallel and Distributed Processing Symposium (IPDPS)*, May 2020.
- [79] J. Wang, T. Zhang, J. Song, N. Sebe, and H. T. Shen, "A survey on learning to hash," *IEEE Transactions on Pattern Analysis and Machine Intelligence*, vol. 40, 2018.

# Supplementary Material for BLEND: A Fast, Memory-Efficient, and Accurate Mechanism to Find Fuzzy Seed Matches

## S1. Generating the Same Hash Value for Similar Seeds

Our goal is to enable generating the same hash values for highly similar seeds. To this end, BLEND uses the SimHash technique [1, 2] to generate a hash value for a seed  $S_i$  using its k-mers  $K_{S_i}$ , as explained in the main paper. Here, we explain in three steps how two hash values,  $B(S_k)$  and  $B(S_l)$ , that BLEND generates respectively for two seeds  $S_k$  and  $S_l$  can have the same value (i.e.,  $B(S_k) = B(S_l)$ ), although  $S_k$  and  $S_l$  may not necessarily be the same seed sequences (i.e.,  $S_k \neq S_l$ ) as stated by the following Theorem S1.

**Theorem S1** *Let  $S_k$  and  $S_l$  be two seed sequences, and  $B(S_k)$  and  $B(S_l)$  be their hash values that BLEND generates, respectively. Then,  $B(S_k) = B(S_l)$  if  $K_{S_k} \simeq K_{S_l}$ .*

First, we assume that the following conditions are true:

1. BLEND uses a *deterministic* hash function to *always* generate the same set of hash values,  $h(K_{S_i})$ , given a set of k-mers  $K_{S_i}$  of a seed  $S_i$ .
2. BLEND *always* generates the hash value  $B(S_i)$  for the given set of hash values of k-mers,  $h(K_{S_i})$ .
3. All the other relevant parameters used in BLEND are the same when generating seeds to ensure  $h(K_{S_k}) = h(K_{S_l})$  if  $S_k = S_l$ . These parameters mainly are k-mer size  $k$ , window length  $w$ , hash function  $h$ , number of neighbor k-mers included in both sets  $n$ , and the seeding mechanism (i.e., BLEND-I or BLEND-S).

Then, we state the following lemma:

**Lemma S1.1**  $B(S_k) = B(S_l)$  if  $S_k = S_l$

To show that Lemma S1.1 is correct, we should show BLEND *always* generates the hash value  $B(S_k)$  for the given set of hash values of k-mers,  $h(K_{S_k})$ . The proof is trivial because the vector encoding step (explained in the main paper) generates the same vectors for each hash value in  $h(K_{S_k})$ , and the vector additions are deterministic such that the addition always generates the same counter vector  $C(S_k)$  given  $h(K_{S_k})$ . The decoding step is also simple and deterministic conversion from a vector to the binary representation of a hash value as explained in the main paper. Thus, BLEND always generates the same  $B(S_k)$  from  $h(K_{S_k})$ .

Lemma S1.1 is true because BLEND always generates the same hash value for  $S_k$  and if  $S_k = S_l$ , then  $h(K_{S_k}) = h(K_{S_l})$ .  $B(S_k) = B(S_l)$  if  $h(K_{S_k}) = h(K_{S_l})$ .

Second, we state the following lemma:

**Lemma S1.2**  $B(S_k)$  and  $B(S_l)$  can be equal (i.e.,  $B(S_k) = B(S_l)$ ) even if two seeds are highly similar but not equal (i.e.,  $S_k \sim S_l$ ).

To prove Lemma S1.2, let us start with a single seed,  $S_k$ , its k-mers,  $K_{S_k}$ , and its hash value after applying the SimHash technique,  $B(S_k)$ . We state the following lemma:

**Lemma S1.3** *Each value in the counter vector (i.e.,  $C(S_k)$ ) used to generate the hash value  $B(S_k)$  can only change by +2 or -2 if we replace a single k-mer in  $K_{S_k}$  with another k-mer (i.e., removing one k-mer and adding another k-mer).*

Lemma S1.3 can easily be proven because a single k-mer is encoded into a single vector as shown in Figure 4 from the main paper, and each position of a vector can only have the value either 1 or -1. A value in the counter vector  $C(S_k)$  at any position  $i$  (i.e.,  $C(S_k)[i]$ ) can only change by +1 or -1 after adding or removing a single vector, respectively. Thus, Lemma S1.3 is true because making two changes (i.e., replacing a k-mer) can only change the values in the counter vector by *at most* +2 or -2.

Given that Lemma S1.3 is true, let us replace a single k-mer in  $K_{S_k}$  with another k-mer to make the seed  $S_k$  a different seed:  $S_l$ . Then, we know that each value in the counter vectors of these seeds,  $C(S_k)$  and  $C(S_l)$ , can only differ by *at most* +2 or -2 from Lemma S1.3.

We will now generalize Lemma S1.3 to changing a single character in  $S_k$  rather than changing a single k-mer in  $K_{S_k}$ . Although it is not relevant to consider the character changes in seeds when using the BLEND-S seeding technique as it links multiple gapped k-mers (i.e., a single character change is unlikely to affect the other linked k-mer), it is more important for the BLEND-I seeding technique as it uses immediately overlapping  $n$ -many k-mers to generate a seed. Let us change a single character in  $S_k$  to make a different seed  $S_l$ . Since  $S_k$  is generated from immediately overlapping k-mers, a single character change in  $S_k$  can change *at most*  $k$ -many k-mers in  $K_{S_k}$  and, thus,  $k$ -many hash values in  $h(K_{S_k})$ . Since we know that a single k-mer change can cause at most +2 or -2 difference between the values in  $C(S_k)$  and  $C(S_l)$  at the same positions, there can be at most  $2k$  difference between these values when we change  $k$ -many k-mers. Then, the hash values  $B(S_k)$  and  $B(S_l)$  must be equal if the following condition holds true:

- All the values in  $C(S_k)$  are either greater than  $2k$  or less than  $-2k+1$  (i.e.,  $C(S_k)[i] > 2k$  or  $C(S_k)[i] \leq -2k \forall i$ )

Given the above condition, we will prove  $B(S_k) = B(S_l)$ . Let us assume that  $B(S_k) \neq B(S_l)$ . This means that *at least one* position in both  $C(S_k)$  and  $C(S_l)$  should be one of the following as these are the only conditions that BLEND checks when setting the bits in  $B(S_k)$  and  $B(S_l)$ :

1.  $C(S_k)[i] > 0$  and  $C(S_l)[i] \leq 0$
2.  $C(S_k)[i] \leq 0$  and  $C(S_l)[i] > 0$

However, neither of the conditions above can hold true because:

1. If  $C(S_k)[i] > 0$  then  $C(S_k)[i] > 2k$  and  $C(S_k)[i] - 2k > 0$  thus  $C(S_l)[i] > 0$  (because there can be at most  $2k$  difference between  $C(S_k)[i]$  and  $C(S_l)[i]$  when we change a single character).
2. If  $C(S_k)[i] \leq 0$  then  $C(S_k)[i] \leq -2k$  and  $C(S_k)[i] + 2k \leq 0$  thus  $C(S_l)[i] \leq 0$ .

Thus,  $B(S_k) = B(S_l)$  by contradiction.

We note that we set our conditions (e.g.,  $C(S_k)$  are either greater than  $2k$  or less than  $-2k+1$ ) for the *most extreme cases* of a single character change because we assume the following most extreme cases:

1. We use BLEND-I.
2.  $K_{S_k}$  and  $K_{S_l}$  differ by  $k$ -many k-mers due to a single character change between  $S_k$  and  $S_l$ .
3. If we replace  $k$ -many k-mers of seed  $S_k$  due to the single character change between  $S_k$  and  $S_l$ , the XNOR operation between the hash values of these replaced k-mers returns 0 such that for any replaced k-mer  $k_i \in K_{S_k}$  and  $k_i \notin K_{S_l}$  there is a replacing k-mer  $k_n \in K_{S_l}$  and  $k_n \notin K_{S_k}$  that provides  $h(k_i) \odot h(k_n) = 0$ . This implies that the difference between two encoded vectors of these hash values are *always* 2 at every position because of the opposite bits that these hash values have at each position.

Case 1 above assumes that we use the BLEND-I technique that may have a bigger impact in the final value of the counter vector compared to using the BLEND-S technique when there is a single character change in seed. Case 2 assumes the worst case in the number of k-mers that need to be replaced with a single character change in a seed. This may not necessarily be true if the change of character is at the beginning or towards the end of the seed such that  $k$ -many k-mers do not cover this change. Case 2 also becomes less impactful when seeds include many k-mers such that changing  $k$ -many k-mers have slightly minor effect in the counter vector. Case 3 above is the most extreme case because it is unlikely that the replacing hash values can *always* satisfy the condition  $h(k_i) \odot h(k_n) = 0$ . Since Case 3 is unlikely to always hold true, we can also claim that BLEND is likely to generate the same hash values such that  $B(S_k) = B(S_l)$  if  $S_k$  and  $S_l$  differs by a single character and *some* of the counter vectors in  $C(S_k)$  are *close* to either  $2k$  or  $-2k+1$ .

Given that Lemma S1.1 and Lemma S1.2 are true, then Theorem S1 is true. We conclude that BLEND finds *all* exactly matching seeds while it can generate the same hash value for highly similar seeds.

### S1.1. A Real Example for Generating the Hash Values of Seeds $S_k$ and $S_l$

Our goal is to show how the k-mer length  $k$  and the number of k-mers to include in a seed,  $n$ , affect the final hash value. To this end, we use the following two seeds as found in the Yeast reference genome:  $S_k$  : CGGATGCTACAGTATATACCA and  $S_l$  : ATGCTACAGTATATACCATCT. Both seeds are 21-character long. We use two different parameter settings when generating the hash values of these seeds. The first setting uses  $k = 7$  as the k-mer length and  $n = 15$  as the number of immediately overlapping k-mers to include in a seed so that we can generate the 21-character long seeds  $S_l$  and  $S_k$ . The second setting uses  $k = 15$  as the k-mer length and  $n = 7$  as the number of k-mers to include in a seed. We use the hash64 hash function as provided in the minimap2 implementation to generate the hash values of the k-mers of seeds.

In Supplementary Tables S1 - S8 we show k-mers, the hash values of the k-mers in their binary form, and the gradual change in the counter vectors used to calculate the hash values for seeds  $S_k$  and  $S_l$ . We update the counter vectors based on the bits in the hash values of each k-mer. Finally, we show the hash values of  $S_k$  and  $S_l$  in the last rows of each table. In Supplementary Tables S1- S4, we use  $k = 7$  as the k-mer length and  $n = 15$  as the number of immediately overlapping k-mers to include in a seed. In Supplementary Tables S5- S8, we use  $k = 15$  as the k-mer length and  $n = 7$  as the number of k-mers to include in a seed.

We make two key observations. First, we observe that the hash values of  $S_k$  and  $S_l$  are equal ( $B(S_k) = B(S_l) = 0b11000100011011001110100110110100$ ) when we use a short k-mer with high number of neighbors even though these two seeds differ by 3 k-mers. Second, the hash values of these two seeds are not equal when we use fewer neighbors with larger k-mers. For  $S_k$  we find the hash value  $B(S_k) = 0b011010000100000010111010011000000$  and for  $S_l$  we find  $B(S_l) = 0b00101101101101100001111110001010011$ . We note that the bit positions with large values in their corresponding counter vectors are less likely to differ between two seeds when the seeds have a large number of k-mers in common. This motivates as to design more intelligent hash functions that are aware of the values in the counter vectors to increase the chance of generating the hash value for similar seeds.



**Table S1: Hash Values of the k-mers of seed  $S_k$ : CGGATGCTACAGTATATACCA for  $k = 7$  and  $n = 15$ . We show the most significant 16 bits of the counter vector  $C(S_k)$ . Last row shows the most significant 16 bits of the hash value of the seed.**

K-mer	Hash Value	C[31]	C[30]	C[29]	C[28]	C[27]	C[26]	C[25]	C[24]	C[23]	C[22]	C[21]	C[20]	C[19]	C[18]	C[17]	C[16]
CGGATGC	0b 10100000 01111111 10000110 10110101	1	-1	1	-1	-1	-1	-1	-1	-1	1	1	1	1	1	1	1
GGATGCT	0b 10101101 11110000 01110100 11010000	2	-2	2	-2	0	0	-2	0	0	2	2	2	0	0	0	0
GATGCTA	0b 01000010 01001011 11011001 10011011	1	-1	1	-3	-1	-1	-1	-1	-1	3	1	1	1	-1	1	1
ATGCTAC	0b 11001100 01110101 01010011 00100110	2	0	0	-4	0	0	-2	-2	-2	4	2	2	0	0	0	2
TGCTACA	0b 10110001 01101010 10101001 10100111	3	-1	1	-3	-1	-1	-3	-1	-3	5	3	1	1	-1	1	1
GCTACAG	0b 11000101 11010111 11010101 00100101	4	0	0	-4	-2	0	-4	0	-2	6	2	2	0	0	2	2
CTACAGT	0b 11001001 01111101 01001010 10110101	5	1	-1	-5	-1	-1	-5	1	-3	7	3	3	1	1	1	3
TACAGTA	0b 00101011 01101111 11111000 11111000	4	0	0	-6	0	-2	-4	2	-4	8	4	2	2	2	2	4
ACAGTAT	0b 11100100 01001110 01110101 00011010	5	1	1	-7	-1	-1	-5	1	-5	9	3	1	3	3	3	3
CAGTATA	0b 10010010 00001101 00100011 10110100	6	0	0	-6	-2	-2	-4	0	-6	8	2	0	4	4	2	4
AGTATAT	0b 01110100 00110000 10101100 00000000	5	1	1	-5	-3	-1	-5	-1	-7	7	3	1	3	3	1	3
GTATATA	0b 11001111 10110000 11001001 10010110	6	2	0	-6	-2	0	-4	0	-6	6	4	2	2	2	0	2
TATATAC	0b 10000001 00001000 00101111 01111111	7	1	-1	-7	-3	-1	-5	1	-7	5	3	1	3	1	-1	1
ATATACC	0b 11001100 11100000 00101000 11011010	8	2	-2	-8	-2	0	-6	0	-6	6	4	0	2	0	-2	0
TATACCA	0b 00110100 00000100 11110100 10010100	7	1	-1	-7	-3	1	-7	-1	-7	5	3	-1	1	1	-3	-1
		B[31]	B[30]	B[29]	B[28]	B[27]	B[26]	B[25]	B[24]	B[23]	B[22]	B[21]	B[20]	B[19]	B[18]	B[17]	B[16]
		1	1	0	0	0	1	0	0	0	1	1	0	1	1	0	0

Best results are highlighted with **bold** text.

**Table S2: Hash Values of the k-mers of seed  $S_k$ : CGGATGCTACAGTATATACCA for  $k = 7$  and  $n = 15$ . We show the least significant 16 bits of the counter vector  $C(S_k)$ . Last row shows the least significant 16 bits of the hash value of the seed.**

K-mer	Hash Value	C[15]	C[14]	C[13]	C[12]	C[11]	C[10]	C[9]	C[8]	C[7]	C[6]	C[5]	C[4]	C[3]	C[2]	C[1]	C[0]
CGGATGC	0b 10100000 01111111 10000110 10110101	1	-1	-1	-1	-1	1	1	-1	1	-1	1	1	-1	1	-1	1
GGATGCT	0b 10101101 11110000 01110100 11010000	0	0	0	0	-2	2	0	-2	2	0	0	2	-2	0	-2	0
GATGCTA	0b 01000010 01001011 11011001 10011011	1	1	-1	1	-1	1	-1	-1	3	-1	-1	3	-1	-1	-1	1
ATGCTAC	0b 11001100 01110101 01010011 00100110	0	2	-2	2	-2	0	0	0	2	-2	0	2	-2	0	0	0
TGCTACA	0b 10110001 01101010 10101001 10100111	1	1	-1	1	-1	-1	-1	1	3	-3	1	1	-3	1	1	1
GCTACAG	0b 11000101 11010111 11010101 00100101	2	2	-2	2	-2	0	-2	2	2	-4	2	0	-4	2	0	2
CTACAGT	0b 11001001 01111101 01001010 10110101	1	3	-3	1	-1	-1	-1	1	3	-5	3	1	-5	3	-1	3
TACAGTA	0b 00101011 01101111 11111000 11111000	2	4	-2	2	0	-2	-2	0	4	-4	4	2	-4	2	-2	2
ACAGTAT	0b 11100100 01001110 01110101 00011010	1	5	-1	3	-1	-1	-3	1	3	-5	3	3	-3	1	-1	1
CAGTATA	0b 10010010 00001101 00100011 10110100	0	4	0	2	-2	-2	-2	2	4	-6	4	4	-4	2	-2	0
AGTATAT	0b 01110100 00110000 10101100 00000000	1	3	1	1	-1	-1	-3	1	3	-7	3	3	-5	1	-3	-1
GTATATA	0b 11001111 10110000 11001001 10010110	2	4	0	0	0	-2	-4	2	4	-8	2	4	-6	2	-2	-2
TATATAC	0b 10000001 00001000 00101111 01111111	1	3	1	-1	1	-1	-3	3	3	-7	3	5	-5	3	-1	-1
ATATACC	0b 11001100 11100000 00101000 11011010	0	2	2	-2	2	-2	-4	2	4	-6	2	6	-4	2	0	-2
TATACCA	0b 00110100 00000100 11110100 10010100	1	3	3	-1	1	-1	-5	1	5	-7	1	7	-5	3	-1	-3
		B[15]	B[14]	B[13]	B[12]	B[11]	B[10]	B[9]	B[8]	B[7]	B[6]	B[5]	B[4]	B[3]	B[2]	B[1]	B[0]
		1	1	1	0	1	0	0	1	1	0	1	1	0	1	0	0

Best results are highlighted with **bold** text.

**Table S3: Hash Values of the k-mers of seed  $S_l$ : ATGCTACAGTATATACCATCT for  $k = 7$  and  $n = 15$ . We show the most significant 16 bits of the counter vector  $C(S_l)$ . Last row shows the most significant 16 bits of the hash value of the seed.**

K-mer	Hash Value	C[31]	C[30]	C[29]	C[28]	C[27]	C[26]	C[25]	C[24]	C[23]	C[22]	C[21]	C[20]	C[19]	C[18]	C[17]	C[16]
ATGCTAC	0b 11001100 01110101 01010011 00100110	1	1	-1	-1	1	1	-1	-1	-1	1	1	1	-1	1	-1	1
TGCTACA	0b 10110001 01101010 10101001 10100111	2	0	0	0	0	0	-2	0	-2	2	2	0	0	0	0	0
GCTACAG	0b 11000101 11010111 11010101 00100101	3	1	-1	-1	-1	1	-3	1	-1	3	1	1	-1	1	1	1
CTACAGT	0b 11001001 01111101 01001010 10110101	4	2	-2	-2	0	0	-4	2	-2	4	2	2	0	2	0	2
TACAGTA	0b 00101011 01101111 11111000 11111000	3	1	-1	-3	1	-1	-3	3	-3	5	3	1	1	3	1	3
ACAGTAT	0b 11100100 01001110 01110101 00011010	4	2	0	-4	0	0	-4	2	-4	6	2	0	2	4	2	2
CAGTATA	0b 10010010 00001101 00100011 10110100	5	1	-1	-3	-1	-1	-3	1	-5	5	1	-1	3	5	1	3
AGTATAT	0b 01110100 00110000 10101100 00000000	4	2	0	-2	-2	0	-4	0	-6	4	2	0	2	4	0	2
GTATATA	0b 11001111 10110000 11001001 10010110	5	3	-1	-3	-1	1	-3	1	-5	3	3	1	1	3	-1	1
TATATAC	0b 10000001 00001000 00101111 01111111	6	2	-2	-4	-2	0	-4	2	-6	2	2	0	2	2	-2	0
ATATACC	0b 11001100 11100000 00101000 11011010	7	3	-3	-5	-1	1	-5	1	-5	3	3	-1	1	1	-3	-1
TATACCA	0b 00110100 00000100 11110100 10010100	6	2	-2	-4	-2	2	-6	0	-6	2	2	-2	0	2	-4	-2
ATACCAT	0b 00000111 10111111 11010101 01001100	5	1	-3	-5	-3	3	-5	1	-5	1	3	-1	1	3	-3	-1
TACCATC	0b 01010110 11100111 00100010 11001101	4	2	-4	-4	-4	4	-4	0	-4	2	4	-2	0	4	-2	0
ACCATCT	0b 00010010 11001000 11001010 11100111	3	1	-5	-3	-5	3	-3	-1	-3	3	3	-3	1	3	-3	-1
		B[31]	B[30]	B[29]	B[28]	B[27]	B[26]	B[25]	B[24]	B[23]	B[22]	B[21]	B[20]	B[19]	B[18]	B[17]	B[16]
		1	1	0	0	0	1	0	0	0	1	1	0	1	1	0	0

Best results are highlighted with **bold** text.

**Table S4: Hash Values of the k-mers of seed  $S_I$ : ATGCTACAGTATATACCATCT for  $k = 7$  and  $n = 15$ . We show the least significant 16 bits of the counter vector  $C(S_I)$ . Last row shows the least significant 16 bits of the hash value of the seed.**

K-mer	Hash Value	C[15]	C[14]	C[13]	C[12]	C[11]	C[10]	C[9]	C[8]	C[7]	C[6]	C[5]	C[4]	C[3]	C[2]	C[1]	C[0]
ATGCTAC	0b 11001100 01110101 01010011 00100110	-1	1	-1	1	-1	-1	1	1	-1	-1	1	-1	-1	1	1	-1
TGCTACA	0b 10110001 01101010 10101001 10100111	0	0	0	0	0	-2	0	2	0	-2	2	-2	-2	2	2	0
GCTACAG	0b 11000101 11010111 11010101 00100101	1	1	-1	1	-1	-1	-1	3	-1	-3	3	-3	-3	3	1	1
CTACAGT	0b 11001001 01111101 01001010 10110101	0	2	-2	0	0	-2	0	2	0	-4	4	-2	-4	4	0	2
TACAGTA	0b 00101011 01101111 11111000 11111000	1	3	-1	1	1	-3	-1	1	1	-3	5	-1	-3	3	-1	1
ACAGTAT	0b 11100100 01001110 01110101 00011010	0	4	0	2	0	-2	-2	2	0	-4	4	0	-2	2	0	0
CAGTATA	0b 10010010 00001101 00100011 10110100	-1	3	1	1	-1	-3	-1	3	1	-5	5	1	-3	3	-1	-1
AGTATAT	0b 01110100 00110000 10101100 00000000	0	2	2	0	0	-2	-2	2	0	-6	4	0	-4	2	-2	-2
GTATATA	0b 11001111 10110000 11001001 10010110	1	3	1	-1	1	-3	-3	3	1	-7	3	1	-5	3	-1	-3
TATATAC	0b 10000001 00001000 00101111 01111111	0	2	2	-2	2	-2	-2	4	0	-6	4	2	-4	4	0	-2
ATATACC	0b 11001100 11100000 00101000 11011010	-1	1	3	-3	3	-3	-3	3	1	-5	3	3	-3	3	1	-3
TATACCA	0b 00110100 00000100 11110100 10010100	0	2	4	-2	2	-2	-4	2	2	-6	2	4	-4	4	0	-4
ATACCAT	0b 00000111 10111111 11010101 01001100	1	3	3	-1	1	-1	-5	3	1	-5	1	3	-3	5	-1	-5
TACCATC	0b 01010110 11100111 00100010 11001101	0	2	4	-2	0	-2	-4	2	2	-4	0	2	-2	6	-2	-4
ACCATCT	0b 00010010 11001000 11001010 11100111	1	3	3	-3	1	-3	-3	1	3	-3	1	1	-3	7	-1	-3
		B[15]	B[14]	B[13]	B[12]	B[11]	B[10]	B[9]	B[8]	B[7]	B[6]	B[5]	B[4]	B[3]	B[2]	B[1]	B[0]
		1	1	1	0	1	0	0	1	1	0	1	1	0	1	0	0

Best results are highlighted with **bold** text.

**Table S5: Hash Values of the k-mers of seed  $S_k$ : CGGATGCTACAGTATATACCA for  $k = 15$  and  $n = 7$ . We show the most significant 16 bits of the counter vector  $C(S_k)$ . Last row shows the most significant 16 bits of the hash value of the seed.**

K-mer	Hash Value	C[31]	C[30]	C[29]	C[28]	C[27]	C[26]	C[25]	C[24]	C[23]	C[22]	C[21]	C[20]	C[19]	C[18]	C[17]	C[16]
CGGATGCTACAGTAT	0b 01001010 11101011 00100110 11001101	-1	1	-1	-1	1	-1	1	-1	1	1	1	-1	1	-1	1	1
GGATGCTACAGTATA	0b 01101100 01000011 11111000 11000000	-2	2	0	-2	2	0	0	-2	0	2	0	-2	0	-2	2	2
GATGCTACAGTATAT	0b 01011000 01000101 00110001 11011000	-3	3	-1	-1	3	-1	-1	-3	-1	3	-1	-3	-1	-1	1	3
ATGCTACAGTATATA	0b 11100001 01110100 01100010 01000010	-2	4	0	-2	2	-2	-2	-2	-2	4	0	-2	-2	0	0	2
TGCTACAGTATATAC	0b 10111100 10011010 00111111 01011011	-1	3	1	-1	3	-1	-3	-3	-1	3	-1	-1	-1	-1	1	1
GCTACAGTATATACC	0b 11101010 01001100 01000100 11100001	0	2	2	-2	4	-2	-2	-4	-2	4	-2	-2	0	0	0	0
CTACAGTATATACCA	0b 00101001 10010001 11111100 01010000	-1	1	3	-3	5	-3	-3	-3	-1	3	-3	-1	-1	-1	-1	1
Seed CGGATGCTACAGTATATACCA		B[31]	B[30]	B[29]	B[28]	B[27]	B[26]	B[25]	B[24]	B[23]	B[22]	B[21]	B[20]	B[19]	B[18]	B[17]	B[16]
		0	1	1	0	1	0	0	0	0	1	0	0	0	0	0	1

Best results are highlighted with **bold** text.

**Table S6: Hash Values of the k-mers of seed  $S_k$ : CGGATGCTACAGTATATACCA for  $k = 15$  and  $n = 7$ . We show the least significant 16 bits of the counter vector  $C(S_k)$ . Last row shows the least significant 16 bits of the hash value of the seed.**

K-mer	Hash Value	C[15]	C[14]	C[13]	C[12]	C[11]	C[10]	C[9]	C[8]	C[7]	C[6]	C[5]	C[4]	C[3]	C[2]	C[1]	C[0]
CGGATGCTACAGTAT	0b 01001010 11101011 00100110 11001101	-1	-1	1	-1	-1	1	1	-1	1	1	-1	-1	1	1	-1	1
GGATGCTACAGTATA	0b 01101100 01000011 11111000 11000000	0	0	2	0	0	0	0	-2	2	2	-2	-2	0	0	-2	0
GATGCTACAGTATAT	0b 01011000 01000101 00110001 11011000	-1	-1	3	1	-1	-1	-1	3	3	3	-3	-1	1	-1	-3	-1
ATGCTACAGTATATA	0b 11100001 01110100 01100010 01000010	-2	0	4	0	-2	-2	0	-2	2	4	-4	-2	0	-2	-2	-2
TGCTACAGTATATAC	0b 10111100 10011010 00111111 01011011	-3	-1	5	1	-1	-1	1	-1	1	5	-5	-1	1	-3	-1	-1
GCTACAGTATATACC	0b 11101010 01001100 01000100 11100001	-4	0	4	0	-2	0	0	-2	2	6	-4	-2	0	-4	-2	0
CTACAGTATATACCA	0b 00101001 10010001 11111100 01010000	-3	1	5	1	-1	1	-1	-3	1	7	-5	-1	-1	-5	-3	-1
Seed CGGATGCTACAGTATATACCA		B[15]	B[14]	B[13]	B[12]	B[11]	B[10]	B[9]	B[8]	B[7]	B[6]	B[5]	B[4]	B[3]	B[2]	B[1]	B[0]
		0	1	1	1	0	1	0	0	1	1	0	0	0	0	0	0

Best results are highlighted with **bold** text.

**Table S7: Hash Values of the k-mers of seed  $S_I$ : ATGCTACAGTATATACCATCT for  $k = 15$  and  $n = 7$ . We show the most significant 16 bits of the counter vector  $C(S_I)$ . Last row shows the most significant 16 bits of the hash value of the seed.**

K-mer	Hash Value	C[31]	C[30]	C[29]	C[28]	C[27]	C[26]	C[25]	C[24]	C[23]	C[22]	C[21]	C[20]	C[19]	C[18]	C[17]	C[16]
ATGCTACAGTATATA	0b 11100001 01110100 01100010 01000010	1	1	1	-1	-1	-1	-1	1	-1	1	1	1	-1	1	-1	-1
TGCTACAGTATATAC	0b 10111100 10011010 00111111 01011011	2	0	2	0	0	0	-2	0	0	0	0	2	0	0	0	-2
GCTACAGTATATACC	0b 11101010 01001100 01000100 11100001	3	1	3	-1	1	-1	-1	-1	-1	1	-1	1	1	1	-1	-3
CTACAGTATATACCA	0b 00101001 10010001 11111100 01010000	2	0	4	-2	2	-2	-2	0	0	0	-2	2	0	0	-2	-2
TACAGTATATACCAT	0b 00001110 00100000 11011100 11110110	1	-1	3	-3	3	-1	-1	-1	-1	-1	-1	1	-1	-1	-3	-3
ACAGTATATACCATC	0b 00101111 10110101 00010000 11011111	0	-2	4	-4	4	0	0	0	0	-2	0	2	0	-2	-2	-4
CAGTATATACCATCT	0b 01100101 11100111 10111011 00111011	-1	-1	5	-5	5	1	-1	1	1	-1	1	1	-1	-1	-1	-3
Seed ATGCTACAGTATATACCATCT		B[31]	B[30]	B[29]	B[28]	B[27]	B[26]	B[25]	B[24]	B[23]	B[22]	B[21]	B[20]	B[19]	B[18]	B[17]	B[16]
		0	0	1	0	1	1	0	1	1	0	1	1	0	0	0	0

Best results are highlighted with **bold** text.

**Table S8: Hash Values of the k-mers of seed  $S_I$ : ATGCTACAGTATATACCATCT for  $k = 15$  and  $n = 7$ . We show the least significant 16 bits of the counter vector  $C(S_I)$ . Last row shows the least significant 16 bits of the hash value of the seed.**

K-mer	Hash Value				C[15]	C[14]	C[13]	C[12]	C[11]	C[10]	C[9]	C[8]	C[7]	C[6]	C[5]	C[4]	C[3]	C[2]	C[1]	C[0]
ATGCTACAGTATATA	0b	11100001	01110100	01100010	01000010	-1	1	1	-1	-1	-1	1	-1	-1	1	-1	-1	-1	-1	-1
TGCTACAGTATATAC	0b	10111100	10011010	00111111	01011011	-2	0	2	0	0	2	0	-2	2	-2	0	0	-2	2	0
GCTACAGTATATACC	0b	11101010	01001100	01000100	11100001	-3	1	1	-1	-1	1	1	-1	-1	3	-1	-1	-1	-3	1
CTACAGTATATACCA	0b	00101001	10010001	11111100	01010000	-2	2	2	0	0	2	0	-2	-2	4	-2	0	-2	-4	0
TACAGTATATACCAT	0b	00001110	00100000	11011100	11110110	-1	3	1	1	1	3	-1	-3	-1	5	-1	1	-3	-3	1
ACAGTATATACCATC	0b	00101111	10111010	00010000	11011111	-2	2	0	2	0	2	-2	-4	0	6	-2	2	-2	-2	2
CAGTATATACCATCT	0b	01100101	11100111	10111011	00111011	-1	1	1	3	1	1	-1	-3	-1	5	-1	3	-1	-3	3
Seed ATGCTACAGTATATACCATCT					B[15]	B[14]	B[13]	B[12]	B[11]	B[10]	B[9]	B[8]	B[7]	B[6]	B[5]	B[4]	B[3]	B[2]	B[1]	B[0]
					0	1	1	1	1	1	0	0	0	1	0	1	0	0	1	1

Best results are highlighted with **bold** text.



## S2. Parameter Exploration

### S2.1. The trade-off between BLEND-I and BLEND-S

Our goal is to show the performance and accuracy trade-offs between the seeding techniques that BLEND supports: BLEND-I and BLEND-S. In Supplementary Figures S1 and S2, we show the performance and peak memory usage comparisons when using BLEND-I and BLEND-S as the seeding technique by keeping all the other relevant parameters identical (e.g., number of k-mers to include in a seed  $n$ , window length  $w$ ). In Supplementary Table S9 we show the assembly quality comparisons in terms of the accuracy and contiguity of the assemblies that we generate using the overlaps that BLEND-I and BLEND-S find. In Supplementary Tables S10 and S11 we show the read mapping quality and accuracy results using these two seeding techniques, respectively.

We also show the values for different parameters we test with BLEND in Supplementary Table S12. We determine the default parameters of BLEND empirically based on the combination of best performance, memory overhead, and accuracy results.




Figure S1: CPU time and peak memory footprint comparisons of read overlapping.




Figure S2: CPU time and peak memory footprint comparisons of read mapping.

Table S9: Assembly quality comparisons between BLEND-I and BLEND-S.

Dataset	Tool	Average Identity (%)	Genome Fraction (%)	K-mer Compl. (%)	Aligned Length (Mbp)	Misassembly Ratio (%)	NGA50 (Kbp)	Average GC (%)	Assembly Length (Mbp)	Largest Contig (Mbp)	NG50 (Kbp)
<i>Human CHM13</i>	BLEND-S	<b>99.8526</b>	<b>98.4847</b>	<b>90.15</b>	<b>3,092.59</b>	<b>0.0108</b>	5,442	40.78	<b>3,095.210</b>	22.840	5,442
	BLEND-I	99.7578	96.4306	84.07	2,990.41	0.0503	<b>8,802</b>	<b>40.84</b>	2,994.976	<b>41.834</b>	<b>9,005</b>
	Reference	100	100	100	3,054.83	0	154,260	40.85	3,054.832	248.387	154,260
<i>D. ananassae</i>	BLEND-S	<b>99.7856</b>	<b>97.2308</b>	<b>86.43</b>	240.39	0.1230	<b>792</b>	41.75	<b>247.153</b>	<b>6.233</b>	<b>799</b>
	BLEND-I	99.7088	97.0399	79.43	<b>248.99</b>	<b>0.1020</b>	363	<b>41.86</b>	258.483	4.439	363
	Reference	100	100	100	213.81	0	26,427	41.81	213.818	30.673	26,427
<i>E. coli</i>	BLEND-S	<b>99.8320</b>	99.8801	<b>87.91</b>	<b>5.12</b>	<b>0.0344</b>	3,417	50.53	5.122	3,417	3,417
	BLEND-I	99.7166	<b>99.8824</b>	80.37	5.04	0.9810	<b>4,025</b>	<b>50.52</b>	<b>5.043</b>	<b>4.946</b>	<b>4,946</b>
	Reference	100	100	100	5.05	0	4,945	50.52	5.046	4.946	4,945
<i>Yeast (PacBio)</i>	BLEND-S	<b>90.3347</b>	83.8814	<b>18.47</b>	<b>0.558</b>	N/A	N/A	<b>38.71</b>	22.952	0.265	116
	BLEND-I	89.1582	<b>97.6297</b>	11.13	0.227	N/A	N/A	38.80	<b>13.679</b>	<b>1.105</b>	<b>551</b>
	Reference	100	100	100	12.16	0	924	38.15	12.157	1.532	924
<i>Yeast (ONT)</i>	BLEND-S	<b>91.0865</b>	7.9798	1.57	0.066	N/A	N/A	<b>38.35</b>	0.900	0.043	N/A
	BLEND-I	89.7622	<b>99.2982</b>	<b>13.68</b>	<b>0.377</b>	N/A	N/A	38.66	<b>12.164</b>	<b>1.554</b>	<b>825</b>
	Reference	100	100	100	12.16	0	924	38.15	12.157	1.532	924

Best results are highlighted with **bold** text. For most metrics best results are the ones closest to the corresponding value of the reference genome.

The best results for *Aligned Length* are determined by the highest number within each dataset. We do not highlight the reference results as the best results.

N/A indicates that we could not generate the corresponding result because either the tool failed or QUAST failed to generate the statistic.

Table S10: Read mapping quality comparisons between BLEND-I and BLEND-S.

Dataset	Tool	Mean Depth of Cov. (×)	Breadth of Coverage (%)	Aligned Reads (#)	Properly Paired (%)
<i>Human CHM13</i>	BLEND-S	16.58	99.99	3,171,916	NA
	BLEND-I	16.58	99.99	<b>3,172,313</b>	NA
<i>D. ananassae</i>	BLEND-S	57.37	<b>99.66</b>	1,223,388	NA
	BLEND-I	<b>57.50</b>	99.65	<b>1,249,731</b>	NA
<i>E. coli</i>	BLEND-S	99.14	99.90	39,048	NA
	BLEND-I	99.14	99.90	<b>39,065</b>	NA
<i>Yeast (PacBio)</i>	BLEND-S	170.43	99.98	217,975	NA
	BLEND-I	<b>195.84</b>	99.98	<b>269,804</b>	NA
<i>Yeast (ONT)</i>	BLEND-S	75.74	99.93	95,847	NA
	BLEND-I	<b>97.86</b>	<b>99.97</b>	<b>134,721</b>	NA
<i>Yeast (Illumina)</i>	BLEND-S	79.63	99.97	6,466,782	95.28
	BLEND-I	<b>79.93</b>	99.97	<b>6,494,489</b>	<b>95.89</b>

Best results are highlighted with **bold** text.

Properly paired rate is only available for paired-end Illumina reads.

Table S11: Read mapping accuracy comparisons between BLEND-I and BLEND-S.

Dataset	Tool	Error Rate (%)	High Quality True Mappings (#)
<i>Yeast (PacBio)</i>	BLEND-S	0.4027985	206,243
	BLEND-I	<b>0.2524054</b>	<b>266,382</b>
<i>Yeast (ONT)</i>	BLEND-S	0.3328221	85,081
	BLEND-I	<b>0.2404970</b>	<b>133,228</b>

Best results are highlighted with **bold** text.

Table S12: Performance, memory, and accuracy comparisons using different parameter settings in BLEND.

Tool	K-mer Length ( $k$ )	# of k-mers in a Seed ( $n$ )	Window Length ( $w$ )	CPU Time (seconds)	Peak Memory (KB)	Average Identity (%)	Genome Fraction (%)
BLEND	9	11	200	62.38	1,115,384	99.7255	99.8502
BLEND	9	13	200	58.13	994,120	99.7294	99.7808
BLEND	9	15	200	49.79	1,030,148	99.7411	99.7619
BLEND	9	17	200	45.03	960,080	99.7302	99.7460
BLEND	9	21	200	36.84	976,456	99.7257	99.6640
BLEND	15	5	200	83.05	1,168,612	99.6735	99.7625
BLEND	15	7	200	74.93	1,137,360	99.7009	99.5874
BLEND	15	11	200	58.09	1,051,912	99.7149	99.1166
BLEND	19	5	200	77.16	1,130,604	99.7312	99.8802
BLEND	19	7	200	50.50	1,078,596	99.7880	99.8424
BLEND	19	11	200	46.26	977,060	99.8078	99.6438
BLEND	21	5	200	67.85	1,116,684	99.7472	99.8835
BLEND	21	7	200	61.63	1,042,724	99.7969	99.8605
BLEND	21	11	200	42.35	969,184	99.8340	99.7515
BLEND	25	5	200	65.61	1,057,804	99.7769	99.8818
BLEND	25	7	200	54.88	1,029,888	99.8320	99.8801
BLEND	25	11	200	37.01	936,260	99.8646	99.8001
BLEND	25	15	200	29.83	866,208	99.8838	99.7307
BLEND	25	17	200	29.59	826,456	99.8784	99.7521
BLEND	25	21	200	26.09	791,736	99.8774	99.6955
BLEND	9	11	50	263.82	1,786,516	99.7013	99.8612
BLEND	9	13	50	411.24	1,805,800	99.6995	99.8573
BLEND	9	15	50	271.00	1,729,784	99.6798	99.8517
BLEND	9	17	50	238.52	1,690,912	99.6690	99.8083
BLEND	9	21	50	206.76	1,725,168	99.6496	99.8150
BLEND	15	5	50	330.84	1,785,456	99.6634	99.8604
BLEND	15	7	50	337.95	1,812,052	99.6280	99.8177
BLEND	15	11	50	236.82	1,803,816	99.5831	99.6893
BLEND	19	5	50	328.67	1,692,248	99.7077	99.8794
BLEND	19	7	50	295.57	1,713,940	99.7188	99.8579
BLEND	19	11	50	201.79	1,700,412	99.7015	99.8578
BLEND	21	5	50	378.58	1,625,388	99.7120	99.8832
BLEND	21	7	50	278.56	1,695,476	99.7333	99.8832
BLEND	21	11	50	189.33	1,694,820	99.7623	99.8594
BLEND	25	5	50	323.69	1,685,304	99.7272	99.8831
BLEND	25	7	50	211.78	1,647,984	99.7722	99.8831
BLEND	25	11	50	170.60	1,683,736	99.8094	99.8866
BLEND	25	15	50	142.42	1,622,452	99.8170	99.8576
BLEND	25	17	50	103.96	1,590,776	99.8073	99.8206
BLEND	25	21	50	109.62	1,548,228	99.7792	99.7880
BLEND	9	11	20	837.50	2,769,552	99.6916	99.8784
BLEND	9	13	20	813.50	2,765,480	99.6834	99.8785
BLEND	9	15	20	764.91	2,795,848	99.6797	99.8756
BLEND	9	17	20	739.52	2,823,188	99.6801	99.8802

We use the *E.coli* dataset for all these runs

## S2.2. The trade-off between BLEND and minimap2

Our goal is to compare BLEND and minimap2 using the same set of parameters that BLEND uses when generating its results. To achieve this, we control the following two conditions. First, we ensure that we use the same seeding technique that minimap2 uses. To this end, we use the BLEND-I seeding technique, which finds minimizer k-mers and uses them as seeds, similar to how minimap2 generates the seeds. We should note that BLEND-I does not always provide the best results in terms of performance or accuracy for the HiFi reads as the default seeding technique is BLEND-S for HiFi datasets in BLEND.

Second, we use the same seed length when we compare BLEND with minimap2. In minimap2, the seed length equals to the k-mer length as minimap2 directly hashes these k-mers and finds the minimizer k-mers from these hash values. However, the seed length in BLEND-I is determined by both the k-mer length and the number of k-mers that we include in a seed. For example, BLEND uses the BLEND-I seeding technique with the k-mer length  $k = 7$  and the number of neighbors  $n = 15$  for the PacBio CLR reads. Combining immediately overlapping 15-many 7-mers generates seeds with length  $15 + 7 - 1 = 21$ . Thus, BLEND-I uses seeds of length 21 while allowing fuzzy seeds to match with its hashing technique. We calculate the seed lengths that BLEND-I uses with its default parameters for each dataset: PacBio HiFi, PacBio CLR, ONT, and Illumina short reads in read overlapping and read mapping. We apply the same seed length and the window length that BLEND uses to minimap2 using the  $k$  and  $w$  parameters. We show these parameters in Supplementary Tables S17 and S18 in the Minimap-Eq rows. In the results we show below, Minimap-Eq indicates the runs of minimap2 when using the same set of parameters that BLEND uses with the BLEND-I seeding technique. We should note that BLEND uses seeds of length 31 (i.e.,  $k = 25$  and  $n = 7$ ) for HiFi datasets in read overlapping while it is not possible to set a seed length larger than 28 in minimap2 due to the limitations in its implementation. Thus, for these datasets, we cannot perform the same experiment with minimap2.

In Supplementary Figures S3 and S4, we show the performance and peak memory usage comparisons when using BLEND with the BLEND-I seeding technique, minimap2, and minimap2-Eq. In Supplementary Table S13 we show the assembly quality comparisons in terms of the accuracy and contiguity of the assemblies that we generate using the overlaps that each tool finds. In Supplementary Tables S14 and S15 we show the read mapping quality and accuracy results, respectively.

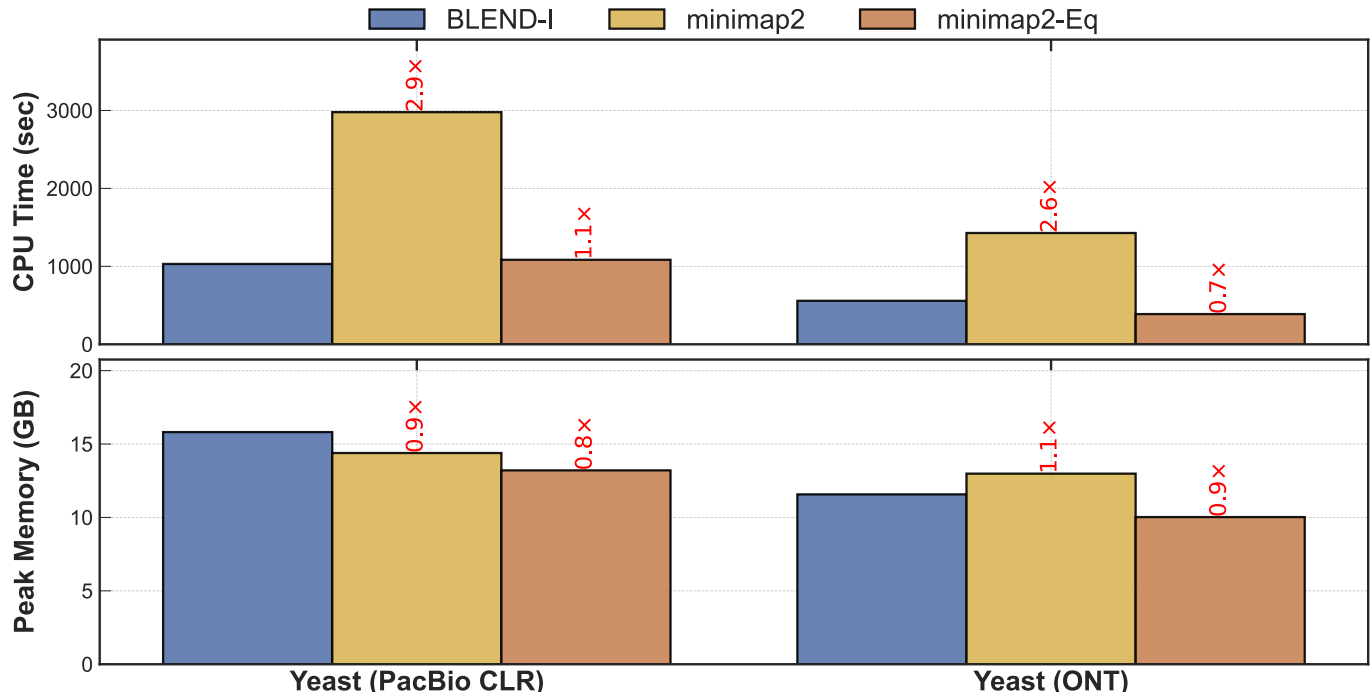


Figure S3: CPU time and peak memory footprint comparisons of read overlapping.






Figure S4: CPU time and peak memory footprint comparisons of read mapping.

Table S13: Assembly quality comparisons when using the parameters equivalent to BLEND-I.

Dataset	Tool	Average Identity (%)	Genome Fraction (%)	K-mer Compl. (%)	Aligned Length (Mbp)	Misassembly Ratio (%)	NGA50 (Kbp)	Average GC (%)	Assembly Length (Mbp)	Largest Contig (Mbp)	NG50 (Kbp)
Yeast (PacBio)	BLEND-I	89.1582	<b>97.6297</b>	<b>11.13</b>	<b>0.227</b>	N/A	N/A	<b>38.80</b>	13.679	1.105	551
	minimap2	88.9002	96.9709	9.74	0.195	N/A	N/A	38.85	<b>12.333</b>	<b>1.561</b>	<b>828</b>
	minimap2-Eq	<b>89.2166</b>	97.2674	10.48	0.215	N/A	N/A	38.82	12.424	1.534	781
	Reference	100	100	100	12.16	0	924	38.15	12.157	1.532	924
Yeast (ONT)	BLEND-I	89.7622	99.2982	<b>13.68</b>	<b>0.377</b>	N/A	N/A	38.66	<b>12.164</b>	1.554	825
	minimap2	88.9393	<b>99.6878</b>	12.06	0.328	N/A	N/A	38.74	12.373	<b>1.560</b>	<b>942</b>
	minimap2-Eq	<b>89.7653</b>	97.3273	13.39	0.371	N/A	N/A	<b>38.64</b>	11.828	1.073	677
	Reference	100	100	100	12.16	0	924	38.15	12.157	1.532	924

Best results are highlighted with **bold** text. For most metrics best results are the ones closest to the corresponding value of the reference genome.

The best results for *Aligned Length* are determined by the highest number within each dataset. We do not highlight the reference results as the best results.

N/A indicates that we could not generate the corresponding result because either the tool failed or QUAST failed to generate the statistic.

Table S14: Read mapping quality comparisons when using the parameters equivalent to BLEND-I.

Dataset	Tool	Mean Depth of Cov. ( $\times$ )	Breadth of Coverage (%)	Aligned Reads (#)	Properly Paired (%)
<i>Human CHM13</i>	BLEND-I	16.58	99.99	<b>3,172,313</b>	NA
	minimap2	16.58	99.99	3,172,261	NA
	minimap2-Eq	16.58	99.99	3,172,312	NA
<i>D. ananassae</i>	BLEND-I	57.50	99.65	1,249,731	NA
	minimap2	<b>57.57</b>	<b>99.67</b>	1,245,931	NA
	minimap2-Eq	57.49	99.65	<b>1,250,816</b>	NA
<i>E. coli</i>	BLEND-I	99.14	99.90	39,065	NA
	minimap2	99.14	99.90	39,065	NA
	minimap2-Eq	99.14	99.90	39,065	NA
<i>Yeast (PacBio)</i>	BLEND-I	195.84	99.98	269,804	NA
	minimap2	<b>195.86</b>	99.98	<b>269,935</b>	NA
	minimap2-Eq	195.81	99.98	269,493	NA
<i>Yeast (ONT)</i>	BLEND-I	97.86	<b>99.97</b>	134,721	NA
	minimap2	<b>97.88</b>	99.96	<b>134,885</b>	NA
	minimap2-Eq	97.82	99.96	134,578	NA
<i>Yeast (Illumina)</i>	BLEND-I	<b>79.93</b>	99.97	<b>6,494,489</b>	<b>95.89</b>
	minimap2	79.91	99.97	6,492,994	<b>95.89</b>
	minimap2-Eq	79.83	99.97	6,485,540	95.72

Best results are highlighted with **bold** text.

Properly paired rate is only available for paired-end Illumina reads.

Table S15: Read mapping accuracy comparisons when using the parameters equivalent to BLEND-I.

Dataset	Tool	Error Rate (%)	High Quality True Mappings (#)
<i>Yeast (PacBio)</i>	BLEND-I	0.2524054	<b>266,382</b>
	minimap2	0.2504307	265,709
	Minimap-Eq	<b>0.2293195</b>	266,355
<i>Yeast (ONT)</i>	BLEND-I	<b>0.2404970</b>	133,228
	minimap2	0.2468770	<b>133,784</b>
	Minimap-Eq	0.2489263	132,641

Best results are highlighted with **bold** text.

### S3. GC Content Distribution for Assembly Quality Assessment

In Figure S5, we show the GC content distribution of assemblies from each dataset that we generate using the overlapping reads from BLEND, minimap2 [3], and MHAP [4].




Figure S5: GC content distribution of assemblies.

## S4. Read Mapping Accuracy

In Supplementary Table S16, we show the accuracy results of each read mapping tool as provided by the `paftools mapeval`, which is a script included in the GitHub page `minimap2` [3]: <https://github.com/lh3/minimap2/tree/master/misc>. In order to generate these accuracy results, we first use `PBSIM2` [5] to generate the simulated reads along with their true mapping locations. Then, the `paftools pbsim2fq` tool takes the true mapping information from `PBSIM2` and generates the reads such that the read IDs in the generated FASTA file are annotated with the true mapping information. Then, `paftools mapeval` uses this annotated information to measure the read mapping accuracy by comparing the mapping result that each tools generates with the true mapping location. For more detailed information we refer to the GitHub page of `paftools` provided above.

**Table S16: Read mapping accuracy comparisons.**

Dataset	Tool	Error Rate (%)	High Quality True Mappings (#)
<i>Yeast</i> (PacBio)	BLEND	0.2524054	<b>266,382</b>
	minimap2	0.2504307	265,709
	LRA	99.9958863	11
	Winnowmap2	<b>0.2474206</b>	247,024
	S-conLSH	97.2181306	3,822
<i>Yeast</i> (ONT)	BLEND	<b>0.2404970</b>	133,228
	minimap2	0.2468770	<b>133,784</b>
	LRA	99.9954840	5
	Winnowmap2	0.2534777	126,899
	S-conLSH	96.6753628	2,361

Best results are highlighted with **bold** text.

## S5. Parameters and Versions

In Supplementary Table S17, we show the parameters we use with `BLEND`, `minimap2`, and `MHAP` [4] for read overlapping. Since there are no default parameters for `minimap2` and `MHAP` when using the HiFi reads, we used the parameters as suggested by the `HiCanu` tool [6]. We found these parameters in the source code of `Canu`. For `minimap2` and `MHAP`, the HiFi parameters are found in the GitHub pages<sup>1,2</sup>, respectively. In Supplementary Table S18, we show the parameters we use with `BLEND`, `minimap2` [3], `LRA` [7], `Winnowmap2` [8, 9], and `S-conLSH` [10, 11] for read mapping. In both Supplementary Tables S17 and S18, *minimap-Eq* shows the parameters that are equivalent to the parameters we use with `BLEND` without the fuzzy seed matching capability. In Supplementary Table S19, we show the version numbers of each tool we use. When calculating the performance and peak memory usage of each tool we use the `time` command from Linux and append the following command to the beginning of each of our runs: `/usr/bin/time -vp`.

<sup>1</sup><https://github.com/marbl/canu/blob/404540a944664cfab00617f4f4fa37be451b34e0/src/pipelines/canu/OverlapMMap.pm#L63-L65>

<sup>2</sup><https://github.com/marbl/canu/blob/404540a944664cfab00617f4f4fa37be451b34e0/src/pipelines/canu/OverlapMhap.pm#L100-L131>

**Table S17: Parameters we use in our evaluation for each tool and dataset in read overlapping.**

<b>Tool</b>	<b>Dataset</b>	<b>Parameters</b>
BLEND	<i>Human CHM13</i>	-x ava-hifi -t 32 (ava-hifi: -strobemers -k25 -w200 -neighbors 7)
BLEND	<i>D. ananassae</i>	-x ava-hifi -t 32
BLEND	<i>E. coli</i>	-x ava-hifi -t 32
BLEND	<i>Yeast PacBio CLR</i>	-x ava-pb -t 32 (ava-pb: -Hk19 -w10 -neighbors 5)
BLEND	<i>Yeast PacBio ONT</i>	-x ava-ont -t 32 (ava-ont: -k15 -w10 -neighbors 5)
minimap2	<i>Human CHM13</i>	-x ava-pb -Hk21 -w14 -t 32
minimap2	<i>D. ananassae</i>	-x ava-pb -Hk21 -w14 -t 32
minimap2	<i>E. coli</i>	-x ava-pb -Hk21 -w14 -t 32
minimap2	<i>Yeast PacBio CLR</i>	-x ava-pb -t 32
minimap2	<i>Yeast PacBio ONT</i>	-x ava-ont -t 32
minimap2-Eq	<i>Yeast PacBio CLR</i>	-x ava-pb -k23 -w10 -t 32
minimap2-Eq	<i>Yeast PacBio ONT</i>	-x ava-ont -k19 -w10 -t 32
MHAP	<i>Human CHM13</i>	-store-full-id -ordered-kmer-size 18 -num-hashes 128 -num-min-matches 5 -ordered-sketch-size 1000 -threshold 0.95 -num-threads 32
MHAP	<i>D. ananassae</i>	-store-full-id -ordered-kmer-size 18 -num-hashes 128 -num-min-matches 5 -ordered-sketch-size 1000 -threshold 0.95 -num-threads 32
MHAP	<i>E. coli</i>	-store-full-id -ordered-kmer-size 18 -num-hashes 128 -num-min-matches 5 -ordered-sketch-size 1000 -threshold 0.95 -num-threads 32
MHAP	<i>Yeast PacBio CLR</i>	-store-full-id -num-threads 32
MHAP	<i>Yeast PacBio ONT</i>	-store-full-id -num-threads 32



**Table S18: Parameters we use in our evaluation for each tool and dataset in read mapping.**

Tool	Dataset	Parameters
BLEND	<i>Human CHM13</i>	-ax map-hifi -t 32 -secondary=no (map-hifi: -strobemers -k19 -w50 -neighbors 5)
BLEND	<i>D. ananassae</i>	-ax map-hifi -t 32 -secondary=no
BLEND	<i>E. coli</i>	-ax map-hifi -t 32 -secondary=no
BLEND	<i>Yeast PacBio CLR</i>	-ax map-pb -t 32 -secondary=no (map-pb: -Hk7 -w10 -neighbors 15)
BLEND	<i>Yeast PacBio ONT</i>	-ax map-ont -t 32 -secondary=no (map-ont: -k7 -w10 -neighbors 11)
BLEND	<i>Yeast Illumina</i>	-ax sr -t 32
minimap2	<i>Human CHM13</i>	-ax map-hifi -t 32 -secondary=no
minimap2	<i>D. ananassae</i>	-ax map-hifi -t 32 -secondary=no
minimap2	<i>E. coli</i>	-ax map-hifi -t 32 -secondary=no
minimap2	<i>Yeast PacBio CLR</i>	-ax map-pb -t 32 -secondary=no
minimap2	<i>Yeast PacBio ONT</i>	-ax map-ont -t 32 -secondary=no
minimap2	<i>Yeast Illumina</i>	-ax sr -t 32
minimap2-Eq	<i>Human CHM13</i>	-ax map-hifi -k23 -w50 -t 32 -secondary=no
minimap2-Eq	<i>D. ananassae</i>	-ax map-hifi -k23 -w50 -t 32 -secondary=no
minimap2-Eq	<i>E. coli</i>	-ax map-hifi -k23 -w50 -t 32 -secondary=no
minimap2-Eq	<i>Yeast PacBio CLR</i>	-ax map-pb -k21 -w10 -t 32 -secondary=no
minimap2-Eq	<i>Yeast PacBio ONT</i>	-ax map-ont -k17 -w10 -t 32 -secondary=no
minimap2-Eq	<i>Yeast Illumina</i>	-ax sr -k25 -w11 -t 32
Winnowmap2	<i>Human CHM13</i>	meryl count k=15 meryl print greater-than distinct=0.9998 -ax map-pb -t 32
Winnowmap2	<i>D. ananassae</i>	meryl count k=15 meryl print greater-than distinct=0.9998 -ax map-pb -t 32
Winnowmap2	<i>E. coli</i>	meryl count k=15 meryl print greater-than distinct=0.9998 -ax map-pb -t 32
Winnowmap2	<i>Yeast PacBio CLR</i>	meryl count k=15 meryl print greater-than distinct=0.9998 -ax map-pb -t 32
Winnowmap2	<i>Yeast PacBio ONT</i>	meryl count k=15 meryl print greater-than distinct=0.9998 -ax map-pb-clr -t 32
LRA	<i>Human CHM13</i>	align -CCS -t 32 -p s
LRA	<i>D. ananassae</i>	align -CCS -t 32 -p s
LRA	<i>E. coli</i>	align -CCS -t 32 -p s
LRA	<i>Yeast PacBio CLR</i>	align -CLR -t 32 -p s
LRA	<i>Yeast PacBio ONT</i>	align -ONT -t 32 -p s
S-conLSH	<i>Human CHM13</i>	-threads 32 -align 1
S-conLSH	<i>E. coli</i>	-threads 32 -align 1
S-conLSH	<i>Yeast PacBio CLR</i>	-threads 32 -align 1
S-conLSH	<i>Yeast PacBio ONT</i>	-threads 32 -align 1

**Table S19: Versions of each tool.**

<b>Tool</b>	<b>Version</b>	<b>GitHub or Conda Link to the Version</b>
BLEND	1.0	<a href="https://github.com/CMU-SAFARI/BLEND">https://github.com/CMU-SAFARI/BLEND</a>
minimap2	2.24	<a href="https://github.com/lh3/minimap2/releases/tag/v2.24">https://github.com/lh3/minimap2/releases/tag/v2.24</a>
MHAP	2.1.3	<a href="https://anaconda.org/bioconda/mhap/2.1.3/download/noarch/mhap-2.1.3-hdfd78af_1.tar.bz2">https://anaconda.org/bioconda/mhap/2.1.3/download/noarch/mhap-2.1.3-hdfd78af_1.tar.bz2</a>
LRA	1.3.2	<a href="https://anaconda.org/bioconda/lra/1.3.2/download/linux-64/lra-1.3.2-ha140323_0.tar.bz2">https://anaconda.org/bioconda/lra/1.3.2/download/linux-64/lra-1.3.2-ha140323_0.tar.bz2</a>
Winnowmap2	2.03	<a href="https://anaconda.org/bioconda/Winnowmap/2.03/download/linux-64/Winnowmap2-2.03-h2e03b76_0.tar.bz2">https://anaconda.org/bioconda/Winnowmap/2.03/download/linux-64/Winnowmap2-2.03-h2e03b76_0.tar.bz2</a>
S-conLSH	2.0	<a href="https://github.com/anganachakraborty/S-conLSH-2.0/tree/292fbe0405f10b3ab63fc3a86cba2807597b582e">https://github.com/anganachakraborty/S-conLSH-2.0/tree/292fbe0405f10b3ab63fc3a86cba2807597b582e</a>

## Supplementary References

- [1] M. S. Charikar, “Similarity Estimation Techniques from Rounding Algorithms,” in *Proceedings of the Thiry-Fourth Annual ACM Symposium on Theory of Computing*, ser. STOC ’02. New York, NY, USA: Association for Computing Machinery, 2002.
- [2] G. S. Manku, A. Jain, and A. Das Sarma, “Detecting Near-Duplicates for Web Crawling,” in *Proceedings of the 16th International Conference on World Wide Web*, ser. WWW ’07. New York, NY, USA: Association for Computing Machinery, 2007.
- [3] H. Li, “Minimap2: pairwise alignment for nucleotide sequences,” *Bioinformatics*, vol. 34, Sep. 2018.
- [4] K. Berlin, S. Koren, C.-S. Chin, J. P. Drake, J. M. Landolin, and A. M. Phillippy, “Assembling large genomes with single-molecule sequencing and locality-sensitive hashing,” *Nature Biotechnology*, vol. 33, Jun. 2015.
- [5] Y. Ono, K. Asai, and M. Hamada, “PBSIM2: a simulator for long-read sequencers with a novel generative model of quality scores,” *Bioinformatics*, vol. 37, Mar. 2021.
- [6] S. Nurk, B. P. Walenz, A. Rhie, M. R. Vollger, G. A. Logsdon, R. Grothe, K. H. Miga, E. E. Eichler, A. M. Phillippy, and S. Koren, “HiCanu: accurate assembly of segmental duplications, satellites, and allelic variants from high-fidelity long reads,” *bioRxiv*, Jan. 2020.
- [7] J. Ren and M. J. P. Chaisson, “lra: A long read aligner for sequences and contigs,” *PLOS Computational Biology*, vol. 17, Jun. 2021.
- [8] C. Jain, A. Rhie, N. F. Hansen, S. Koren, and A. M. Phillippy, “Long-read mapping to repetitive reference sequences using Winnowmap2,” *Nature Methods*, Apr. 2022.
- [9] C. Jain, A. Rhie, H. Zhang, C. Chu, B. P. Walenz, S. Koren, and A. M. Phillippy, “Weighted minimizer sampling improves long read mapping,” *Bioinformatics*, vol. 36, Jul. 2020.
- [10] A. Chakraborty and S. Bandyopadhyay, “conLSH: Context based Locality Sensitive Hashing for mapping of noisy SMRT reads,” *Computational Biology and Chemistry*, vol. 85, Apr. 2020.
- [11] A. Chakraborty, B. Morgenstern, and S. Bandyopadhyay, “S-conLSH: alignment-free gapped mapping of noisy long reads,” *BMC Bioinformatics*, vol. 22, Feb. 2021.